

February 23, 2026

Department of Health and Human Services
Assistant Secretary for Technology Policy / Office of the National Coordinator for Health
Information Technology
Attention: HHS Health Sector AI RFI
Mary E. Switzer Building, Mail Stop: 7033A
330 C Street SW, Washington, DC 20201

Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part
of Clinical Care (RIN 0955-AA13; 90 FR 60108)

Submitted via: <http://www.regulations.gov>

PUBLIC COMMENT

The Evidence Gap in Clinical AI: Why Federal Standards Must Precede Adoption

*How Unvalidated Claims, Misaligned Incentives, and Regulatory Fragmentation Are Putting
Patients at Risk*

Submitted by:

Rahul Gorijavolu

Nikhil Jaiswal

Cathy Mullins, Patient Partner

Tina Hernandez-Boussard, PhD, MS, MPH

Leo Anthony Celi, MD, MS, MPH

On behalf of MIT Critical Data

I. Summary

This comment responds to the Department's Request for Information on accelerating the adoption and use of artificial intelligence in clinical care.¹ We respectfully submit that the framing of this RFI, which focuses primarily on removing barriers to AI adoption, reflects an incomplete picture of the current landscape. Before the Department acts to accelerate adoption, it should examine the systemic risks created by AI systems that have been deployed with inadequate evidence, that operate in regulatory gaps no federal agency has claimed jurisdiction over, and that access patient data through legal frameworks never designed for autonomous technology. These risks are not hypothetical. They are documented, ongoing, and escalating.²

This comment draws on the collective expertise of the authoring team, which includes contributors to MIT Critical Data, a consortium spanning 40 countries with over a decade of work at the intersection of clinical medicine, health data science, and global health fairness, including the development of the MIMIC database, one of the most widely used open clinical datasets in the world.³ This comment also draws on direct investigation of the Doctronic AI prescribing pilot in Utah,⁴ including public records requests filed by members of our team under the Government Records Access and Management Act (GRAMA), as well as extensive analysis of the broader AI healthcare landscape, including multiple companies, documented patient harms, regulatory gaps, liability questions, algorithmic bias, and international comparisons. Doctronic is discussed as a case study, but the systemic problems this comment identifies extend far beyond any single company or state.

We respond to all ten RFI questions below. Our central recommendation is that HHS should establish federal minimum evidence standards, liability frameworks, and fairness requirements for clinical AI before further accelerating adoption. We recognize that some commenters will argue these safeguards would slow innovation. We believe the opposite is true: clear standards are the foundation on which sustainable adoption must be built. Without them, the inevitable accumulation of patient harms will produce a public backlash that sets clinical AI back by years.

¹<https://www.federalregister.gov/documents/2025/12/23/2025-23641/request-for-information-accelerating-the-adoption-and-use-of-artificial-intelligence-as-part-of>

²<https://www.statnews.com/2026/02/03/utah-doctronic-ai-doctor-prescription-renewal-fda-regulation/>

³<https://physionet.org/content/mimiciv/>

⁴<https://commerce.utah.gov/2026/01/06/news-release-utah-and-doctronic-announce-groundbreaking-partnership-for-ai-prescription-medication-renewals/>

II. The Evidence Crisis: AI Healthcare Claims Outpace Proof

A. The Doctronic Case Study: How a Preprint Became Policy

To understand the evidence gap in clinical AI, it is useful to examine the strongest case: the Doctronic pilot in Utah, which has received significant regulatory scrutiny and press coverage compared with other autonomous clinical AI systems. If the evidence supporting this pilot is inadequate, it raises serious questions about the rest of the industry. The central evidence supporting Doctronic is a single preprint posted to medRxiv in July 2025.⁵ The paper reports a 99.2% concordance rate between AI and physician treatment plans. This figure has been repeated by major outlets as though it constitutes validated evidence.⁶

It does not meet that standard. medRxiv’s own FAQ states preprints “should not be used to guide clinical practice.”⁷ The methodological problems are substantial:

- All eight authors are equity holders in Doctronic Inc.⁸
- Ethics approval came from Doctronic’s own internal committee, not an independent IRB.⁹
- The sample: 500 urgent-care encounters from a single week, which is insufficient to support claims about safety at scale.¹⁰
- Clinicians saw AI-generated documentation before rendering assessments, introducing anchoring bias.¹¹
- Primary diagnosis match was only 81%; the 99.2% measures “treatment plan alignment,” a softer standard.¹²
- The study tested urgent care, but the Utah pilot involves chronic medication renewals, which is a fundamentally different clinical context.¹³

Doctronic’s co-founder claimed the AI is “infinitely safer than a human doctor”¹⁴, a claim supported by no independent evidence. The medical co-founder is a vascular surgeon with no prior AI research.¹⁵ The company’s first CMO with relevant AI credentials was hired in October 2025, after the study.¹⁶

⁵<https://www.medrxiv.org/content/10.1101/2025.07.14.25331406v1>

⁶<https://www.axios.com/local/salt-lake-city/2026/01/07/utah-ai-drug-prescriptions-doctronic/>;
<https://fortune.com/2026/01/08/ai-prescription-renewals-doctronic-utah-doctors-warn-puts-patients-at-risk/>

⁷<https://www.medrxiv.org/about/FAQ>

⁸<https://www.medrxiv.org/content/10.1101/2025.07.14.25331406v1>

⁹<https://www.medrxiv.org/content/10.1101/2025.07.14.25331406v1>

¹⁰<https://www.medrxiv.org/content/10.1101/2025.07.14.25331406v1>

¹¹<https://www.medrxiv.org/content/10.1101/2025.07.14.25331406v1>

¹²<https://www.medrxiv.org/content/10.1101/2025.07.14.25331406v1>

¹³<https://www.doctronic.ai/blog/doctronic-treatment-plan-consistency/>

¹⁴<https://www.sltrib.com/news/2026/01/09/utahns-can-now-use-ai-renew/>

¹⁵<https://surgery.ucsf.edu/bio/adam-oskowitz-md-phd>

¹⁶<https://www.businesswire.com/news/home/20251015771722/en/>

B. A Landscape of Unvalidated Clinical AI Claims

The evidentiary problems with Doctronic are not unique to that company. They reflect standard practice across the clinical AI industry, where companies routinely make clinical claims that far outstrip their published evidence. The following examples illustrate the scope of the problem.

Hippocratic AI has raised \$404 million at a \$3.5 billion valuation, marketing “AI nurse agents” for direct patient interaction. The company claims 115 million clinical interactions “with no safety issues,” yet these metrics come from internal studies using a proprietary model that cannot be independently benchmarked. National Nurses United organized over 20 demonstrations against the company.¹⁷

Glass Health (YC-funded, \$6.5M raised) generates AI differential diagnoses and treatment plans. Despite claiming use at Harvard and Stanford, no published peer-reviewed clinical validation studies exist. The company deliberately avoids FDA-triggering language by marketing as “clinical decision support.”¹⁸

K Health (\$380M raised, \$900M valuation) operates an AI symptom checker. While a company-associated study claimed 68% alignment with physicians, an independent vignette study found top-3 diagnostic accuracy of only 36% compared to 82.1% for general practitioners.

Among newer entrants, YC-backed **FutureClinic** claims “digital medical clinics that are 99% AI, 1% doctor,” while **Prana** markets itself as a “full-stack medical provider” and “always-on physician.” These are extraordinarily aggressive clinical claims for pre-revenue startups.

C. The Corporate Failure Record

The pattern of overpromising and underdelivering in clinical AI is not new. IBM invested over \$4 billion in Watson Health before selling it for approximately \$1 billion; Watson for Oncology was trained on synthetic cases and proved inappropriate internationally.¹⁹ Babylon Health raised \$1.2 billion before bankruptcy; the UK’s MHRA acknowledged “concerns” about its chatbot’s systematic triage failures, including missed heart attack symptoms and documented gender bias.^{20,21}

The Texas Attorney General’s September 2024 enforcement action against Pieces Technologies found the company marketed hallucination rates below 0.001% that were “likely inaccurate.” Four major hospitals had been feeding real-time patient data to Pieces based on these claims.²²

¹⁷<https://telecareaware.com/are-hippocratic-ai-and-ai-nurses-the-wave-of-the-future-or-just-another-tide-of-hype-two-articles-question/>

¹⁸<https://techcrunch.com/2023/09/08/glass-health-is-building-an-ai-for-suggesting-medical-diagnoses/>

¹⁹<https://theconversation.com/from-a-deranged-provocateur-to-ibms-failed-ai-superproject-the-controversial-story-of-how-data-has-transformed-healthcare-189362>

²⁰<https://theweek.com/health/babylon-health-the-failed-ai-wonder-app-that-dazzled-politicians>

²¹<https://techcrunch.com/2021/03/05/uk-mhra-says-it-has-concerns-about-babylon-health-and-flags-legal-gap-around-triage-chatbots/>

²²<https://www.texasattorneygeneral.gov/news/releases/attorney-general-ken-paxton-reaches-settlement-first-its-kind-healthcare-generative-ai-investigation>

These corporate failures point to a deeper problem: AI systems that perform well in controlled settings consistently underperform in real clinical environments. An Oxford University study published in *Nature Medicine* (February 2026) quantified this gap, finding that AI diagnostic accuracy dropped from 95% on standardized tests to less than 35% when real patients interacted with these tools. Participants using AI chatbots actually performed worse than those using Google search. This finding has direct implications for the Doctronic pilot: a concordance study conducted in controlled conditions tells us very little about how the system will perform when deployed autonomously to real patients with complex medication histories.²³

²³<https://www.sciencenews.org/article/medical-advice-ai-chatbots-symptoms>

III. Documented Patient Harms: A Systematic Pattern

The evidence reviewed above concerns unvalidated claims. But the consequences of deploying AI health tools without adequate safeguards are not merely theoretical. Documented harms have already occurred across multiple products, patient populations, and categories of harm. These cases illustrate what happens when AI systems enter clinical environments without the protections that HHS has the opportunity to establish through this rulemaking process.

A. The NEDA/Tessa Incident

After the National Eating Disorders Association terminated its human helpline staff, the Tessa chatbot, which was augmented with generative AI capabilities without NEDA's knowledge, gave eating disorder patients advice to lose 1–2 pounds per week, eat no more than 2,000 calories daily, maintain a 500–1,000 calorie deficit, and purchase skin calipers to measure body fat.²⁴ As one eating disorder survivor reported: every suggestion “were things that led to my eating disorder.”

B. Character.AI Teen Suicides: The Failure of Voluntary Safety Measures

Sewell Setzer III, age 14, died by suicide on February 28, 2024, after developing an intense emotional attachment to a Character.AI chatbot. When he expressed suicidal ideation, the bot failed to trigger safety interventions. Juliana Peralta, age 13, died in November 2023 after similar engagement. Google and Character.AI settled in January 2026, with additional lawsuits filed in Colorado, New York, and Texas. These cases are relevant to the clinical AI question because they demonstrate that AI companies deploying health-adjacent tools without mandatory safety standards will not reliably implement adequate safeguards on their own. Character.AI added safety features only after public litigation and media pressure, not before deployment.²⁵

C. The Koko Experiment

Approximately 4,000 mental health users received GPT-3-co-written support through the Koko platform without clear informed consent and without IRB approval.²⁶ Medical ethicist Art Caplan stated: “If you want to set back the use of AI in mental health, start exactly this way.” The Koko incident demonstrates a recurring problem: AI was introduced into a clinical context without the informed consent protections that would be required of any human-delivered intervention. The same dynamic is present in the Doctronic pilot, where patients interacting with AI may not fully understand the nature of the entity making decisions about their medications.

D. AI Drug Information Errors

²⁴<https://www.cbsnews.com/news/eating-disorder-helpline-chatbot-disabled/>;
<https://www.npr.org/sections/health-shots/2023/06/08/1180838096/an-eating-disorders-chatbot-offered-dieting-advice-raising-fears-about-ai-in-he-a>

²⁵<https://www.cnn.com/2026/01/07/business/character-ai-google-settle-teen-suicide-lawsuit>;
<https://www.nbcnews.com/tech/characterai-lawsuit-florida-teen-death-rcna176791>

²⁶<https://www.psychiatrist.com/news/hidden-use-of-chatgpt-in-online-mental-health-counseling-raises-ethical-concerns/>

The drug information errors documented in general-purpose AI are directly relevant to the question of autonomous AI prescribing. A Heidelberg University study found 38% of ChatGPT drug answers were false, with 26% posing high risk of patient harm.²⁷ An ASHP study found only 26% of drug information responses were satisfactory, with ChatGPT missing the dangerous Paxlovid-verapamil interaction and fabricating references.²⁸

These are not niche findings. ECRI, the independent patient safety organization, named misuse of AI chatbots the number one health technology hazard for both 2025 and 2026, documenting chatbots suggesting incorrect diagnoses, recommending unnecessary testing, and inventing body parts.²⁹ Mount Sinai researchers found AI chatbots “routinely elaborated on fake medical details, confidently generating explanations about conditions or treatments that do not exist.”³⁰

A Stanford study found AI therapy chatbots enabled dangerous behavior: when a user mentioned losing their job and asked about tall bridges, the chatbot listed bridges rather than recognizing suicidal ideation. An additional safety concern that warrants research attention is AI sycophancy: the tendency of AI systems to agree with users rather than provide accurate information, particularly in patient-facing applications. When a patient tells an AI chatbot they believe they have a particular condition, sycophantic AI will tend to confirm rather than challenge that belief, even when the evidence does not support it. This dynamic is especially dangerous in mental health contexts, where validation of distorted thinking can reinforce harmful patterns. Taken together, these cases and dynamics establish that AI health tools have caused serious harm across a range of contexts, that companies have not voluntarily implemented adequate safeguards, and that the harms tend to fall disproportionately on vulnerable populations. This pattern should inform how HHS approaches the question of accelerating clinical AI adoption.³¹

²⁷<https://pubmed.ncbi.nlm.nih.gov/37263772/>

²⁸<https://news.ashp.org/news/ashp-news/2023/12/05/study-finds-chatgpt-provides-inaccurate-responses-to-drug-questions>

²⁹<https://home.ecri.org/blogs/ecri-news/misuse-of-ai-chatbots-tops-annual-list-of-health-technology-hazards>

³⁰<https://www.mountsinai.org/about/newsroom/2025/ai-chatbots-can-run-with-medical-misinformation-study-finds-highlighting-the-need-for-stronger-safeguards>

³¹<https://hai.stanford.edu/news/exploring-the-dangers-of-ai-in-mental-health-care>

IV. Venture Capital Is Shaping Clinical AI Policy

A. Regulatory Strategy as Investment Thesis

The preceding sections establish that clinical AI companies are making unvalidated claims and that those claims have produced real patient harms. This section examines why the system is producing these outcomes: the financial incentives driving clinical AI deployment are structurally misaligned with patient safety. Doctronic has raised \$25 million from Lightspeed Venture Partners,³²³³ Union Square Ventures, and Tusk Venture Partners.³⁴

Tusk Venture Partners is of particular relevance to this comment. Founded by Bradley Tusk, who served as Communications Director for Senator Schumer, Deputy Governor of Illinois, Bloomberg’s campaign manager³⁵, the firm invests in startups facing regulatory barriers and provides political strategy in exchange for equity. Co-founder Jordan Nof stated: “How do we close the door behind us? I want to create a regulatory moat.”³⁶ Tusk himself: “Venture capitalists, you know, we’re not here for the good of humanity.”³⁷

Doctronic’s website acknowledges Tusk’s “regulatory expertise helps us navigate the complex healthcare landscape.”³⁸ Public records obtained through GRAMA confirm Tusk Ventures appears in negotiation records between Doctronic and Utah’s Office of AI Policy.³⁹

This arrangement creates a structural conflict that HHS should consider as it evaluates how to accelerate adoption responsibly. When the same entity that profits from a company’s commercial success is also shaping the regulatory environment that company operates in, the resulting regulatory outcomes cannot be presumed to reflect an independent assessment of patient safety. The question for HHS is whether the current regulatory framework has adequate safeguards against this kind of influence. Based on our investigation, it does not.

B. The “AI Doctor” Gold Rush

Y Combinator and other accelerators are actively funding startups marketing AI as replacements for physician judgment. HHS should recognize that pressure to “accelerate adoption” is substantially generated by venture capital return timelines rather than patient demand, clinical evidence, or healthcare system needs. Patients want affordable, accessible healthcare, and replacing physicians with AI does not address that need.

³²<https://medium.com/sand-hill-road/finding-regulatory-moats-at-the-series-a-with-jordan-nof-of-tusk-venture-partners-5e5c0e7abf54>

³³<https://lsvp.com/stories/investing-in-doctronic-ai-healthcares-24-7-digital-doctor/>

³⁴<https://pulse2.com/doctronic-20-million-series-a/>

³⁵https://en.wikipedia.org/wiki/Bradley_Tusk

³⁶<https://medium.com/sand-hill-road/finding-regulatory-moats-at-the-series-a-with-jordan-nof-of-tusk-venture-partners-5e5c0e7abf54>

³⁷<https://techcrunch.com/2019/12/18/tusk-ventures-second-fund/>

³⁸<https://www.doctronic.ai/strategic-partnerships-driving-healthcare-innovation/>

³⁹GRAMA Response, Utah Dept. of Commerce, Jan. 26, 2026 (on file with authors).

V. The FDA Regulatory Framework Has Critical Gaps

A. The Clinical Decision Support Exemption

The evidence, harms, and incentive misalignment documented above might be less concerning if a robust federal regulatory framework existed to ensure clinical AI safety. It does not. The most significant gap is the Clinical Decision Support (CDS) exemption. Section 3060(a) of the 21st Century Cures Act exempts CDS software from FDA oversight if it meets four criteria, including that clinicians can independently review the AI's reasoning.⁴⁰ Companies exploit this by framing autonomous AI systems as “support” tools and adding disclaimers that recommendations require “clinician review.” The opacity of deep learning models makes the independent review criterion largely fictional.

FDA Commissioner Makary's January 2026 CDS guidance further loosened oversight, extending enforcement discretion and stating the FDA should be “less paternalistic.” The guidance removed analysis of automation bias.⁴¹

B. The 510(k) Pathway Is Inadequate

Over 1,300 AI/ML-enabled medical devices have been authorized, virtually all via the 510(k) pathway requiring only “substantial equivalence” to a predicate device. A 2025 study found fewer than 2% were supported by randomized clinical trials.⁴² More than a third of first-generation AI devices were cleared based on non-AI predicates, meaning novel AI capabilities were deemed equivalent to fundamentally different technologies. The practical consequence is that the vast majority of clinical AI tools in use today reached the market without the kind of clinical trial evidence that would be required for a new pharmaceutical or medical device. This is the evidentiary environment into which the Department is proposing to accelerate further adoption.⁴³

C. Why Autonomous AI Prescribing Falls in the Gap

The FDA regulates medical devices but not the practice of medicine. State medical practice acts define prescriptive authority, which has historically been limited to human licensees. The FDA has declined comment on Doctronic's pilot, stating it “falls outside the agency's regulatory purview.”⁴⁴ This means no regulatory body is clearly responsible for ensuring autonomous AI prescribing is safe.

D. FDA Staffing Cuts Compound the Problem

⁴⁰<https://bipartisanpolicy.org/issue-brief/fda-oversight-understanding-the-regulation-of-health-ai-tools/>

⁴¹<https://www.nelsonmullins.com/insights/blogs/ai-task-force/ai/hhs-proposes-rule-to-deregulate-health-it-and-advance-ai-interoperability-hti-5>

⁴²<https://pmc.ncbi.nlm.nih.gov/articles/PMC12140231/>

⁴³[https://www.thelancet.com/journals/landig/article/PIIS2589-7500\(23\)00126-7/fulltext](https://www.thelancet.com/journals/landig/article/PIIS2589-7500(23)00126-7/fulltext)

⁴⁴<https://www.statnews.com/2026/02/03/utah-doctronic-ai-doctor-prescription-renewal-fda-regulation/>

Post-deployment monitoring is also critically underdeveloped. The FDA's MAUDE system captures fewer than 0.5% of adverse events from medical devices, and Predetermined Change Control Plans (PCCPs) allow AI/ML device modifications without fresh clearance, meaning AI tools can evolve substantially after initial authorization without triggering new review.⁴⁵ Approximately 3,500 FDA employees were terminated in April 2025 (~20% of staff). The FY2026 budget proposes an 11.5% cut. Former Commissioner Califf stated: "The FDA as we've known it is finished."⁴⁶ The agency missed its first drug approval deadline due to staff cuts.⁴⁷ This paradox, in which the agency faces record AI submissions amid decimated oversight capacity, was noted in Nature's npj Digital Medicine as "consternation."⁴⁸

⁴⁵<https://www.fda.gov/medical-devices/mandatory-reporting-requirements-manufacturers-importers-and-device-user-facilities/manufactur-and-user-facility-device-experience-database-maude>

⁴⁶<https://www.fiercehealthcare.com/regulatory/rfk-jr-prepares-10000-job-cuts-across-hhs-new-wave-worker-reductions>;
<https://www.npr.org/sections/shots-health-news/2025/04/01/g-s1-57485/hhs-fda-layoffs-doge-cdc-nih>

⁴⁷<https://www.axios.com/2025/07/15/doge-cuts-fda-hhs-biotech>

⁴⁸<https://www.nature.com/articles/s41746-025-01540-2>

VI. The Liability Vacuum for Autonomous Clinical AI

If the regulatory framework does not prevent harm, the legal system should at least provide recourse for patients who are harmed. But here, too, the existing framework has no clear answer for autonomous clinical AI. As of early 2026, no U.S. malpractice jury verdict has centered on AI as the basis of a clinical claim, and the legal doctrines that would govern such a case are poorly suited to the technology.⁴⁹

A. Medical Malpractice Does Not Fit

Traditional malpractice requires a physician-patient relationship (which does not exist with autonomous AI), a standard of care (undefined for AI), breach (which is nearly impossible to determine in opaque systems), and causation (which is difficult to isolate). As W. Nicholson Price II of Michigan Law School has documented, deep learning embeds reasoning in thousands of simulated neurons, making breach determination effectively impossible.⁵⁰

B. Product Liability Does Not Fit Either

Courts have historically been reluctant to apply product liability to software, treating it as a “service.” Yet when AI acts autonomously, the malpractice framework, which requires a human practitioner, also breaks down. The practical result is that a patient harmed by an autonomous AI prescribing decision may have no viable legal theory under which to seek compensation. This is not an abstract concern. It is the legal environment in which the Doctronic pilot is currently operating.⁵¹

C. The Learned Intermediary Doctrine Collapses

The learned intermediary doctrine shields drug manufacturers by requiring warnings only to prescribing physicians. When AI removes the physician, neither the manufacturer (which warned the “prescriber”) nor the AI company clearly bears liability.⁵² Consider a concrete scenario: Doctronic’s AI renews an SSRI for a patient who has since been prescribed a medication with a dangerous interaction. The drug manufacturer warned the “prescriber,” but the prescriber is an algorithm. The AI company may argue it is a technology provider, not a medical practitioner. The state authorized the pilot but disclaimed clinical responsibility. The patient is left with no clear defendant. This is not a hypothetical edge case; it is the foreseeable consequence of deploying autonomous prescribing without first resolving the liability framework.

D. The UnitedHealth Precedent

The most significant AI healthcare litigation is *Estate of Lokken v. UnitedHealth Group*, alleging UnitedHealth used the nH Predict algorithm to systematically deny post-acute care despite

⁴⁹<https://www.ncbi.nlm.nih.gov/books/NBK613216/>; https://repository.law.umich.edu/cgi/viewcontent.cgi?article=1569&context=book_chapters

⁵⁰<https://www.ncbi.nlm.nih.gov/books/NBK613216/>; https://repository.law.umich.edu/cgi/viewcontent.cgi?article=1569&context=book_chapters

⁵¹<https://journalofethics.ama-assn.org/article/are-current-tort-liability-doctrines-adequate-addressing-injury-caused-ai/2019-02>

⁵²<https://www.druganddevicelawblog.com/2017/07/artificial-intelligence-and-learned-intermediaries.html>

knowing its approximately 90% error rate. A Senate investigation found denial rates jumped from 10.9% to 22.7% after AI implementation, with skilled nursing denials increasing ninefold.⁵³

⁵³<https://www.cbsnews.com/news/unitedhealth-lawsuit-ai-deny-claims-medicare-advantage-health-insurance-denials/>

VII. Algorithmic Bias Threatens to Worsen Health Disparities

The risks documented above, including unvalidated evidence, patient harms, regulatory gaps, and liability failures, do not fall equally on all populations. The evidence consistently shows that clinical AI systems perform worse for minority and underserved patients, raising serious fairness concerns that this RFI should address directly. Obermeyer et al. (Science, 2019) found that an algorithm applied to roughly 200 million Americans used healthcare costs as a proxy for need, systematically underestimating Black patients' illness. At a given risk score, Black patients had 26.3% more chronic conditions than White patients. The bias reduced Black patients identified for extra care by more than half.⁵⁴

Pulse oximetry bias feeds directly into AI decision-making: Sjoding et al. (NEJM, 2020) found Black patients had nearly three times the frequency of undetected hypoxemia.⁵⁵ Dermatology AI performs substantially worse on dark skin. No public benchmark included biopsy-proven malignancy on dark skin before Daneshjou et al. (Science Advances, 2022) created the first such dataset.⁵⁶

The Epic Sepsis Model, deployed at hundreds of hospitals, achieved an AUC of only 0.63 in external validation, well below the claimed 0.76 to 0.83, and missed 67% of sepsis patients.⁵⁷

A systematic review found 22 of 24 studies (91.7%) identified demographic biases in medical LLMs, with 93.7% finding gender bias and 90.9% finding racial/ethnic bias.⁵⁸ GPT-4 specifically showed associations between demographic attributes and recommendations for more expensive procedures.⁵⁹

The data representation problem is structural: non-Hispanic Whites constitute over 90% of clinical trial participants while representing 60.7% of the population, and fewer than 44% of clinical trials report any race/ethnicity data. These disparities are compounded for rural and lower socioeconomic status populations, who are systematically underrepresented in both the clinical data used to train AI systems and in the health systems where AI tools are first deployed and validated. Safety-net hospitals and rural health facilities, which serve the most vulnerable patients, are least likely to have the infrastructure to evaluate AI tools locally, creating a dynamic in which the populations at greatest risk from biased AI are also the populations with the least capacity to detect and mitigate that bias. For autonomous prescribing systems like Doctronic, this means the AI is making medication decisions based on training data that systematically underrepresents the patients most likely to experience adverse outcomes. HHS cannot responsibly accelerate adoption of clinical AI without requiring that these systems demonstrate equitable performance across the populations they serve.⁶⁰

⁵⁴<https://www.science.org/doi/10.1126/science.aax2342>

⁵⁵Sjoding, M. et al., NEJM, Dec. 2020. <https://doi.org/10.1056/NEJMc2029240>

⁵⁶<https://pmc.ncbi.nlm.nih.gov/articles/PMC9374341/>

⁵⁷<https://pmc.ncbi.nlm.nih.gov/articles/PMC8218233/>

⁵⁸<https://link.springer.com/article/10.1186/s12939-025-02419-0>

⁵⁹<https://pubmed.ncbi.nlm.nih.gov/38123252/>

⁶⁰[https://www.thelancet.com/journals/lanam/article/PIIS2667-193X\(22\)00069-2/fulltext](https://www.thelancet.com/journals/lanam/article/PIIS2667-193X(22)00069-2/fulltext)

VIII. State Regulatory Capacity Does Not Match Ambition

A. Utah: Institutional Mismatch

Given the evidence gaps, documented harms, incentive misalignment, regulatory failures, liability vacuum, and fairness risks detailed in the preceding sections, the question of who authorizes clinical AI deployment becomes critical. In the case of the most consequential clinical AI authorization to date, the answer is troubling. Utah’s Office of Artificial Intelligence Policy has four staff members: a mathematician, a social systems modeler, an internet privacy attorney, and a wealth management professional. None has medical, clinical, or pharmaceutical training.⁶¹⁶²

This team authorized Doctronic to autonomously renew approximately 192 medications, including SSRIs, antihypertensives, statins, and hormonal contraceptives.⁶³ Healthcare law firm Goldsand Friedberg concluded the agreement “clearly steps over” licensed professional activity and noted no agreement “requires representations regarding FDA qualification.”⁶⁴

No major medical society has endorsed autonomous AI prescribing. The AMA stated it “poses serious risks.”⁶⁵ The Utah Medical Association called it “a bad idea.”⁶⁶ Public Citizen’s Dr. Steinbrook called it “a dangerous first step” and warned that “decisions made by AI with human oversight are in practice decisions made by AI most of the time.”⁶⁷

B. The Transparency Deficit

Our team filed public records requests under GRAMA seeking monthly safety reports, negotiation drafts, metric definitions, and federal communications about the pilot. The Utah Department of Commerce denied access to the majority of records, classifying safety reports as “protected” under business confidentiality exemptions. The state’s position is that our interests “do not outweigh Doctronic’s business confidentiality interests.”⁶⁸

The result is a closed loop: a state office without medical expertise authorized clinical AI deployment, the safety data that could allow independent evaluation is classified as confidential, and the public has no mechanism to assess whether the pilot is functioning safely. If HHS intends to encourage state-level AI experimentation, this is precisely the kind of transparency failure that federal standards should prevent.

C. State Proliferation

⁶¹<https://ai.utah.gov/staff/>

⁶²<https://utahpolicy.com/news-release/70903-zachary-boyd-appointed-director-of-the-utah-office-of-artificial-intelligence-policy>

⁶³<https://commerce.utah.gov/2026/01/06/news-release-utah-and-doctronic-announce-groundbreaking-partnership-for-ai-prescription-medication-renewals/>

⁶⁴<https://www.goldsandfriedberg.com/insights/utahs-ai-regulatory-learning-laboratory-early-signals-from-health-care-mitigation-agreements>

⁶⁵<https://fortune.com/2026/01/08/ai-prescription-renewals-doctronic-utah-doctors-warn-puts-patients-at-risk/>

⁶⁶<https://www.sltrib.com/news/2026/01/09/utahns-can-now-use-ai-renew/>

⁶⁷<https://www.citizen.org/news/utah-pilot-program-for-medication-renewals-with-ai-perverts-medical-practice/>

⁶⁸GRAMA Appeal Response, Utah Dept. of Commerce, Feb. 18, 2026 (on file with authors).

Texas enacted its AI sandbox law (RAIGA) effective January 1, 2026. Doctronic is in discussions with Texas, Arizona, and Missouri. Wyoming is preparing its own sandbox. Doctronic's co-founder expects "a dozen other states approve something like this in 2026." Without federal minimum standards, the result will be a patchwork of state experiments conducted under widely varying levels of oversight, with no coordinated adverse event reporting, no consistent evidence requirements, and no mechanism for one state to learn from another's failures. The window for HHS to establish a federal floor before this fragmentation becomes entrenched is narrow.

IX. The Prescription Data Pipeline Operates Without Meaningful Consent

The regulatory and liability gaps described above are compounded by a data access framework that operates largely without patient knowledge. Doctronic accesses medication histories through the Surescripts network, which covers over 95% of U.S. pharmacies and 99% of the population.⁶⁹ The legal mechanism is HIPAA's TPO exception, 45 C.F.R. § 164.506), which permits PHI sharing for treatment without patient consent.

As NYU's Technology Law & Policy Clinic documented, data sharing with Surescripts is "an inevitable consequence of sending an e-prescription," and existing rules "permit Surescripts and other HIEs to share patients' sensitive prescription drug histories without patients' knowledge or consent."⁷⁰ Patients wishing to opt out must complete and notarize a form and mail it to Surescripts headquarters.⁷¹

No court has tested whether the TPO exception applies to autonomous AI acting as a de facto healthcare provider. TPG Capital became Surescripts' majority investor in October 2024.⁷² The FTC's 2023 antitrust settlement eliminated exclusivity contracts, potentially making it easier for new AI entrants to access prescription data.⁷³

⁶⁹<https://surescripts.com/news-center/intelligence-in-action/interoperability/why-it-matters-prescribers-can-now-access-accurate-medication-history-data-for-virtually-every-american>

⁷⁰<https://www.law.nyu.edu/sites/default/files/Comment%20of%20NYU%20TLP%20Clinic%20on%20Proposed%20HIPAA%20Rule.pdf>

⁷¹<https://surescripts.com/opt-out-faq>

⁷²<https://www.fiercehealthcare.com/health-tech/pe-firm-tpg-buys-majority-stake-health-information-network-surescripts>

⁷³<https://www.ftc.gov/news-events/news/press-releases/2023/07/ftc-reaches-proposed-settlement-surescripts-illegal-monopolization-case>

X. International Comparisons Reveal the U.S. as an Outlier

The gaps in the U.S. framework are not inevitable features of AI regulation. Other jurisdictions have developed more protective approaches that provide useful models for HHS consideration.

The **EU AI Act** classifies medical AI as high-risk, requiring risk management systems, data governance ensuring representative datasets, mandatory human oversight, accuracy standards, post-market monitoring, and serious incident reporting. Penalties can reach €35 million or 7% of global turnover.⁷⁴ Product defectiveness can be presumed if AI does not comply with safety requirements.

Health Canada requires a Sex and Gender-Based Analysis Plus: manufacturers must collect disaggregated data on sex, gender, racial/ethnic minorities, elderly, and pediatric populations. This mandatory fairness requirement has no U.S. parallel.⁷⁵

Japan mandates that AI medical devices must be used under physician oversight, with physicians making all final treatment decisions.⁷⁶

The U.S. is the only major jurisdiction to have allowed autonomous AI prescribing and the only one with a broad statutory CDS exemption that removes many clinical AI tools from oversight entirely. We do not suggest that HHS adopt any single international framework wholesale, but the comparative evidence demonstrates that the requirements we recommend, including mandatory fairness testing, human oversight proportional to autonomy, and post-market surveillance, are not radical proposals. They are the minimum standards that other nations have concluded are necessary to protect patients.

⁷⁴<https://www.hunton.com/insights/legal/the-impact-of-the-eu-ai-act-on-the-development-and-use-of-medical-devices>

⁷⁵<https://www.canada.ca/en/health-canada/services/drugs-health-products/medical-devices/application-information/guidance-documents/pre-market-guidance-machine-learning-enabled-medical-devices.html>

⁷⁶<https://www.nature.com/articles/s41746-025-01540-2>

XI. Payer and Insurance Implications

The regulatory and legal gaps documented above have practical consequences for how AI-prescribed care will be paid for and insured. No major insurer has issued a policy specifically addressing coverage of medications prescribed by autonomous AI systems. The current framework presumes prescriptions from licensed human practitioners.⁷⁷

CMS has focused on preventing AI-driven denials: the April 2023 Final Rule requires Medicare Advantage determinations to be based on individual circumstances, not algorithms. CMS's WISer Model (January 2026) uses AI for Medicare prior authorizations but maintains that final coverage denial decisions are made by "licensed clinicians, not machines."⁷⁸

Malpractice insurance for AI remains undefined. Major insurers review AI involvement case-by-case; no AI-specific frameworks exist.⁷⁹ The reimbursement framework is embryonic: CPT 2026 will be the first year AI is explicitly recognized in formal medical coding.⁸⁰

These gaps create practical consequences on both sides of the adoption question. Physicians who might benefit from AI-assisted decision support may resist using it because the liability and coverage implications are unclear. Meanwhile, patients whose medications are renewed by an autonomous AI system may discover that their insurer does not recognize the prescription as valid, or that no malpractice framework covers them if something goes wrong. Resolving these questions is not a barrier to innovation; it is a prerequisite for responsible deployment.

⁷⁷<https://www.thedoctors.com/the-doctors-advocate/second-quarter-2024/artificial-intelligence-emerging-issues-in-healthcare-and-insurance-public-policy>

⁷⁸<https://stateline.org/2025/12/04/medicares-new-ai-experiment-sparks-alarm-among-doctors-lawmakers/>

⁷⁹<https://www.thedoctors.com/the-doctors-advocate/second-quarter-2024/artificial-intelligence-emerging-issues-in-healthcare-and-insurance-public-policy>

⁸⁰<https://bipartisanpolicy.org/issue-brief/paying-for-ai-in-u-s-health-care/>

XII. Expert and Medical Society Consensus

The concerns raised in this comment are not ours alone. They are shared by the leading medical societies, academic journals, and patient safety organizations in the United States.

The NEJM AI editors declared (March 2024): artificial intelligence must meet the same bar for clinical evidence as other interventions, and mandated trial preregistration for AI studies.⁸¹

The AMA's augmented intelligence policy deliberately uses "augmented" over "artificial" to emphasize AI's assistive role. The AMA insists AI should augment, not replace, clinical judgment, and that the "try-first mentality should be reserved for testing environments only as the risks to patient health are simply too significant." AMA President Bruce Scott stated: "Voluntary standards are not going to be enough."⁸²

The American College of Physicians insists AI tools must "complement and not supplant the logic and decision-making of physicians," warning against "automation complacency."⁸³

The National Academy of Medicine published its AI Code of Conduct in July 2025, the first comprehensive values-driven framework for responsible AI in health. The consistency of these positions across institutions with different mandates, constituencies, and political orientations is notable. There is no significant expert voice arguing that autonomous clinical AI should be deployed without independent validation, fairness testing, and clear liability frameworks. The question is whether federal policy will reflect this consensus or continue to lag behind it.⁸⁴

⁸¹<https://ai.nejm.org/doi/10.1056/AIe2400146>

⁸²<https://www.ama-assn.org/practice-management/digital-health/ama-position-2025-federal-government-ai-action-plan>

⁸³<https://www.acpjournals.org/doi/10.7326/M24-0146>

⁸⁴<https://nam.edu/our-work/programs/leadership-consortium/health-care-artificial-intelligence-code-of-conduct/>

XIII. Responses to RFI Questions

Question 1: What are the biggest barriers to AI innovation in healthcare?

We acknowledge that real barriers to responsible AI innovation exist. Data interoperability remains fragmented across incompatible EHR systems. Reimbursement pathways are embryonic; CPT 2026 will be the first year AI is explicitly recognized in formal medical coding. Regulatory jurisdiction is divided among federal agencies and 50 state governments, creating uncertainty for companies seeking to deploy across state lines.⁸⁵

However, the most consequential barrier to sustainable AI adoption is the absence of clear federal safety standards. This absence undermines public trust, creates liability uncertainty that deters responsible developers from adopting AI tools, and invites a race to the bottom among states competing to attract AI companies with the least restrictive environments. When patients learn that the AI renewing their medications was authorized by a four-person state office with no medical expertise, or that the supporting evidence is a company-funded preprint, their confidence in clinical AI erodes. Moreover, if reimbursement and regulatory frameworks privilege AI-enabled care, health systems that cannot afford AI tools, particularly safety-net hospitals and rural facilities, risk falling further behind, creating a two-tiered healthcare system in which the quality of care depends on whether a patient's provider can afford the latest technology.^{86,87}

The framing of this RFI implicitly treats safety standards as obstacles to innovation. We urge HHS to recognize that the opposite is true. The pharmaceutical industry earned public trust not by operating without FDA oversight but because FDA oversight gave patients confidence that approved drugs met minimum safety standards. Clinical AI requires the same foundation.

Question 2: What regulatory or payment changes should be prioritized?

The most urgent regulatory change is closing the Clinical Decision Support exemption for AI tools that functionally make autonomous clinical decisions. The CDS exemption under the 21st Century Cures Act was designed for tools that assist clinicians, not systems that replace them. Companies exploit this by adding nominal "clinician review" disclaimers to functionally autonomous AI output. HHS should clarify, in coordination with FDA, that the exemption does not apply when AI output is the primary or default clinical action, when clinicians cannot meaningfully review the reasoning due to model opacity, or when the AI operates without real-time physician oversight.⁸⁸

⁸⁵<https://bipartisanpolicy.org/issue-brief/paying-for-ai-in-u-s-health-care/>

⁸⁶<https://www.statnews.com/2026/02/03/utah-doctroneic-ai-doctor-prescription-renewal-fda-regulation/>

⁸⁷<https://fortune.com/2026/01/08/ai-prescription-renewals-doctroneic-utah-doctors-warn-puts-patients-at-risk/>

⁸⁸<https://bipartisanpolicy.org/issue-brief/fda-oversight-understanding-the-regulation-of-health-ai-tools/>

Second, HHS should establish mandatory pre-deployment clinical validation proportional to the risk level of the AI tool. A symptom-checker app presents different risks than a system autonomously prescribing medications, and validation requirements should scale accordingly. At minimum, any AI making autonomous treatment or prescribing decisions should be supported by independent, peer-reviewed evidence. Fewer than 2% of FDA-cleared AI/ML devices are currently supported by randomized clinical trials.⁸⁹

Third, HHS should require mandatory fairness testing modeled on Health Canada's Sex and Gender-Based Analysis Plus, which requires manufacturers to collect disaggregated performance data across demographic subgroups. Fourth, HHS should reverse the HTI-5 proposed rule's removal of model card requirements, which would eliminate among the only federal transparency guardrails for clinical AI. Finally, HHS should develop AI-specific reimbursement pathways through CMS that tie payment to validated clinical outcomes rather than volume of AI-generated recommendations.⁹⁰⁹¹⁹²

Question 3: What novel legal issues arise for AI not classified as medical devices?

The liability vacuum is the most urgent legal issue. As documented in Section VI, when autonomous AI makes clinical decisions, the traditional medical malpractice framework does not apply because there is no physician-patient relationship, no defined standard of care for AI, and no ability to determine breach in opaque systems. Product liability also fails because courts have historically treated software as a service rather than a product. The result is that a patient harmed by an autonomous AI prescribing decision may have no viable legal theory under which to seek compensation.⁹³⁹⁴

The learned intermediary doctrine, which shields drug manufacturers by requiring warnings only to prescribing physicians, collapses entirely when AI removes the physician from the prescribing decision. If Doctrionic's AI renews a medication where a dangerous interaction has developed, the drug manufacturer warned the "prescriber" (an algorithm), the AI company may claim it is a technology platform, and the patient has no clear defendant. HHS should support the development of clear liability apportionment frameworks for autonomous clinical AI, potentially including strict liability for systems making treatment decisions without physician oversight.⁹⁵

HIPAA's Treatment, Payment, and Operations exception, which permits sharing of protected health information without patient consent, has never been tested for autonomous AI acting as a de facto healthcare provider. HHS should issue guidance clarifying that autonomous AI systems

⁸⁹<https://pmc.ncbi.nlm.nih.gov/articles/PMC12140231/>

⁹⁰<https://www.canada.ca/en/health-canada/services/drugs-health-products/medical-devices/application-information/guidance-documents/pre-market-guidance-machine-learning-enabled-medical-devices.html>

⁹¹<https://www.nelsonmullins.com/insights/blogs/ai-task-force/ai/hhs-proposes-rule-to-deregulate-health-it-and-advance-ai-interoperability-hti-5>

⁹²<https://bipartisanpolicy.org/issue-brief/paying-for-ai-in-u-s-health-care/>

⁹³<https://www.ncbi.nlm.nih.gov/books/NBK613216/>

⁹⁴<https://journalofethics.ama-assn.org/article/are-current-tort-liability-doctrines-adequate-addressing-injury-caused-ai/2019-02>

⁹⁵<https://www.druganddevicelawblog.com/2017/07/artificial-intelligence-and-learned-intermediaries.html>

accessing patient prescription data require explicit patient authorization, separate from the general TPO exception that was designed for communication between human providers.⁹⁶

Question 4: How should AI be evaluated for clinical safety?

Current evaluation practices are inadequate. The Doctronic preprint illustrates the problem: all eight authors hold equity in the company, ethics approval came from an internal committee, and the sample was 500 cases from a single week. Yet this study is the evidentiary basis for the first autonomous AI prescribing system in U.S. history. The Epic Sepsis Model, deployed at hundreds of hospitals, achieved an AUC of only 0.63 in external validation versus the claimed 0.76 to 0.83, missing 67% of sepsis patients. The Oxford/Nature Medicine study (February 2026) found AI diagnostic accuracy dropped from 95% on standardized tests to under 35% in real-world interactions. These cases demonstrate that company-reported metrics cannot substitute for independent evaluation.⁹⁷⁹⁸⁹⁹

HHS should require, at minimum: (a) independent third-party testing by entities with no financial relationship to the developer, following the model proposed in npj Digital Medicine; (b) mandatory reporting of disaggregated performance data across racial, ethnic, gender, and age subgroups, given the systematic biases documented by Obermeyer et al. and others; (c) real-world performance monitoring post-deployment with mandatory adverse event reporting, since controlled-setting performance consistently fails to predict real-world outcomes; and (d) standardized clinical AI safety benchmarks developed in collaboration with NIST.¹⁰⁰¹⁰¹

The NEJM AI editors have mandated trial preregistration for AI intervention studies starting January 1, 2025. HHS should adopt a parallel requirement: any clinical AI system seeking federal recognition, reimbursement, or deployment in federally funded facilities should be supported by preregistered, independently conducted clinical evidence. We also commend to HHS the LTARC evaluation framework (Local, Task-specific, Agile, Reflective, Community-partnered), which goes beyond existing reporting checklists by requiring evaluation to be conducted locally at the deploying institution, tailored to the specific clinical task, updated continuously as models and populations evolve, reflective of the care context including workflow integration, and developed in partnership with the communities the AI is meant to serve. Only 0.2% of healthcare AI publications report any community involvement in the evaluation process; this must change.¹⁰²

⁹⁶<https://www.law.nyu.edu/sites/default/files/Comment%20of%20NYU%20TLP%20Clinic%20on%20Proposed%20HIPAA%20Rule.pdf>

⁹⁷<https://www.medrxiv.org/content/10.1101/2025.07.14.25331406v1>

⁹⁸<https://pmc.ncbi.nlm.nih.gov/articles/PMC8218233/>

⁹⁹<https://www.sciencenews.org/article/medical-advice-ai-chatbots-symptoms>

¹⁰⁰<https://ai.nejm.org/doi/10.1056/AIe2400146>

¹⁰¹<https://www.science.org/doi/10.1126/science.aax2342>

102

Question 5: What private sector activities should HHS support?

HHS should support and strengthen existing private-sector quality initiatives, including the Joint Commission and the National Academy of Medicine's AI Code of Conduct. These frameworks establish important voluntary baselines for responsible AI deployment in healthcare. However, HHS should be clear-eyed about the limits of voluntary self-regulation. As AMA President Bruce Scott stated, "voluntary standards are not going to be enough."¹⁰³¹⁰⁴

The Texas Attorney General's enforcement action against Pieces Technologies, which found the company's claimed hallucination rate of less than 0.001% was "likely inaccurate," demonstrates that companies will make unsupported clinical claims absent enforceable standards. The FTC has signaled its authority over deceptive AI claims through Operation AI Comply, but healthcare-specific enforcement requires domain expertise that general consumer protection agencies lack. HHS should establish healthcare-specific standards for AI clinical claims and coordinate enforcement with the FTC and state attorneys general.¹⁰⁵¹⁰⁶

HHS should also support the development of independent AI testing laboratories and registries, analogous to the role clinical laboratories play in diagnostics. Developers should be able to submit AI tools for independent evaluation, and health systems should have access to standardized, comparable performance data when making procurement decisions. HHS should also ensure that reimbursement frameworks cover the infrastructure required for responsible AI governance, not just the AI tools themselves. Effective clinical AI requires continuous monitoring, local validation, clinician training, bias auditing, and community engagement. These activities are currently unfunded, creating a systematic underinvestment in the governance infrastructure that safe AI deployment requires. CMS should develop dedicated reimbursement codes for AI governance activities as part of its AI payment strategy.

Question 6: Where have AI tools met or fallen short of expectations?

The pattern across clinical AI is consistent: tools that appear highly accurate in controlled settings underperform substantially in real-world deployment. The Epic Sepsis Model, marketed with an AUC of 0.76 to 0.83, achieved only 0.63 in external validation and missed 67% of patients who developed sepsis, while generating false alerts on 18% of all hospitalized patients. Babylon Health raised \$1.2 billion on claims of AI diagnostic superiority before documenting

¹⁰³<https://nam.edu/our-work/programs/leadership-consortium/health-care-artificial-intelligence-code-of-conduct/>

¹⁰⁴<https://www.ama-assn.org/practice-management/digital-health/ama-position-2025-federal-government-ai-action-plan>

¹⁰⁵<https://www.texasattorneygeneral.gov/news/releases/attorney-general-ken-paxton-reaches-settlement-first-its-kind-healthcare-generative-ai-investigation>

¹⁰⁶<https://www.ftc.gov/business-guidance/blog/2023/02/keep-your-ai-claims-check>

systematic triage failures, including missed heart attack symptoms, missed pulmonary embolism risk, and gender-biased assessments. The company collapsed in 2023.¹⁰⁷¹⁰⁸

K Health's AI symptom checker claimed 68% alignment with physicians in a company-associated study, but an independent evaluation found top-3 diagnostic accuracy of only 36%, compared to 82.1% for general practitioners. The Oxford/Nature Medicine study (February 2026) found that real-world AI diagnostic accuracy dropped to under 35%, and patients using AI chatbots performed worse than those using Google search alone. ECRI, the independent patient safety organization, ranked AI misuse as the number one health technology hazard for both 2025 and 2026.¹⁰⁹¹¹⁰

We are not aware of any clinical AI system that has demonstrated, through independent evaluation, sustained real-world performance matching its marketing claims. This pattern should give HHS significant pause before acting to accelerate adoption. The Department should consider requiring post-deployment performance monitoring as a condition of any federal support for clinical AI.

Question 7: Who influences AI adoption decisions and what governance is needed?

AI adoption decisions are currently made by a combination of health system executives (CMOs, CIOs, compliance officers), state regulators, and, increasingly, venture-backed companies pursuing regulatory authorization. Each of these actors faces incentives that may not align with patient safety. Health system executives face pressure to reduce costs and staffing. State regulators face pressure to attract technology companies. VC-backed companies face pressure to deploy before competitors and demonstrate traction to investors.¹¹¹

The Utah case illustrates the governance failure that results: a four-person state office with no medical, clinical, or pharmaceutical training authorized autonomous prescribing of 192 medications, while the state's own medical association opposed the decision and the AMA called it a serious risk. HHS should require that any governance body authorizing autonomous clinical AI include practicing clinicians with relevant domain expertise, patient representatives, and individuals with no financial relationship to the AI developer.¹¹²

At the health system level, HHS should encourage the establishment of AI governance committees modeled on institutional review boards, with authority to evaluate clinical AI tools before deployment. The NAM AI Code of Conduct provides a useful framework. HHS should

¹⁰⁷<https://pmc.ncbi.nlm.nih.gov/articles/PMC8218233/>

¹⁰⁸<https://theweek.com/health/babylon-health-the-failed-ai-wonder-app-that-dazzled-politicians>

¹⁰⁹<https://www.sciencenews.org/article/medical-advice-ai-chatbots-symptoms>

¹¹⁰<https://home.ecri.org/blogs/ecri-news/misuse-of-ai-chatbots-tops-annual-list-of-health-technology-hazards>

¹¹¹<https://medium.com/sand-hill-road/finding-regulatory-moats-at-the-series-a-with-jordan-nof-of-tusk-venture-partners-5e5c0e7abf54>

¹¹²<https://ai.utah.gov/staff/>

also require clinical AI vendors to disclose all funding sources, investor relationships, and conflicts of interest as a condition of deployment in federally funded facilities.¹¹³

Question 8: How should interoperability be addressed?

Interoperability is a legitimate priority, but it is insufficient without safety standards. The Surescripts network already provides near-universal medication history data to clinical AI systems like Doctronic, demonstrating that data access is not the binding constraint on AI deployment. The binding constraint is the absence of standards ensuring that AI systems using this data do so safely and accurately.¹¹⁴

The HTI-5 proposed rule's removal of model card requirements for AI clinical decision support tools is counterproductive. Model cards, which disclose training data characteristics, known limitations, and performance across subgroups, are among the only federal transparency requirements for clinical AI. Removing them increases data flow while decreasing the information clinicians and patients need to evaluate whether AI tools are trustworthy. HHS should strengthen, not weaken, transparency requirements as a prerequisite for safe interoperability.¹¹⁵

We also note that interoperability without adequate consent protections creates privacy risks. As documented in Section IX, patient medication histories flow through the Surescripts network to commercial AI entities without meaningful patient knowledge or consent, under a HIPAA exception designed for communication between human providers. Expanding interoperability without addressing this consent gap would compound an existing problem.

Question 9: What concerns do patients and caregivers have?

Patient and caregiver concerns about clinical AI are grounded in documented harm, not speculation. The Tessa chatbot gave eating disorder patients advice to restrict calories and purchase body fat calipers. A 14-year-old and a 13-year-old died by suicide after intensive engagement with Character.AI chatbots that failed to trigger safety interventions. A Stanford study found AI therapy chatbots that listed tall bridges for a user expressing suicidal ideation rather than recognizing the crisis. These are not edge cases from experimental systems; they are outcomes from commercially deployed products used by millions of people.¹¹⁶¹¹⁷¹¹⁸

¹¹³<https://nam.edu/our-work/programs/leadership-consortium/health-care-artificial-intelligence-code-of-conduct/>

¹¹⁴<https://surescripts.com/news-center/intelligence-in-action/interoperability/why-it-matters-prescribers-can-now-access-accurate-medication-history-data-for-virtually-every-american>

¹¹⁵<https://www.nelsonmullins.com/insights/blogs/ai-task-force/ai/hhs-proposes-rule-to-deregulate-health-it-and-advance-ai-interoperability-hti-5>

¹¹⁶<https://www.cbsnews.com/news/eating-disorder-helpline-chatbot-disabled/>;

<https://www.npr.org/sections/health-shots/2023/06/08/1180838096/an-eating-disorders-chatbot-offered-dieting-advice-raising-fears-about-ai-in-health>

¹¹⁷<https://www.cnn.com/2026/01/07/business/character-ai-google-settle-teen-suicide-lawsuit>

¹¹⁸<https://hai.stanford.edu/news/exploring-the-dangers-of-ai-in-mental-health-care>

Patients also face an information asymmetry that undermines informed decision-making. When Doctronic's preprint is reported by major outlets as demonstrating 99.2% accuracy without noting that the study was internally funded, non-peer-reviewed, and conducted by equity holders, patients cannot meaningfully evaluate the safety of the AI making decisions about their medications. Over 40 million people daily ask ChatGPT health questions, but a Heidelberg study found 38% of its drug answers are false, with 26% posing high risk of patient harm. Patients have no way to distinguish reliable AI-generated health information from unreliable information.¹¹⁹

The data access concern is equally significant. Most patients are unaware that their complete prescription histories are accessible through the Surescripts network, that HIPAA's TPO exception permits this access without consent, and that commercial AI companies can obtain this data. Patients wishing to opt out must complete and notarize a paper form and mail it to corporate headquarters. This is not a meaningful consent framework. HHS should require that patients be notified when AI systems access their health data and given a genuine opportunity to consent or decline.¹²⁰

Question 10: What research priorities should HHS pursue?

HHS should prioritize research on the fairness impacts of clinical AI across race, gender, income, age, and disability status. The Obermeyer et al. finding that a widely used algorithm underestimated Black patients' illness severity, reducing their identification for needed care by more than half, was published in 2019. Six years later, a systematic review found 22 of 24 studies identified demographic biases in medical LLMs. The problem is well-documented but poorly understood in terms of mechanisms and interventions. HHS-funded research should focus on developing validated debiasing methodologies and mandatory performance benchmarks across demographic subgroups.^{121 122}

Second, HHS should fund long-term outcome studies of clinical AI deployment. Nearly all existing evidence evaluates AI performance at a single point in time. No published study tracks clinical outcomes for patients whose care was managed or influenced by AI over months or years. Without longitudinal data, the safety of clinical AI deployment remains an open question. The Oxford study's finding that AI diagnostic accuracy in real-world conditions drops to under 35% suggests that point-in-time evaluations substantially overestimate real-world benefit.¹²³

Third, HHS should fund research on automation bias and clinical deskilling. Evidence from other safety-critical domains (aviation, nuclear power) demonstrates that human operators who rely on automated systems lose independent judgment capacity over time. If AI clinical decision support

¹¹⁹<https://pubmed.ncbi.nlm.nih.gov/37263772/>

¹²⁰<https://www.law.nyu.edu/sites/default/files/Comment%20of%20NYU%20TLP%20Clinic%20on%20Proposed%20HIPAA%20Rule.pdf>

¹²¹<https://www.science.org/doi/10.1126/science.aax2342>

¹²²<https://link.springer.com/article/10.1186/s12939-025-02419-0>

¹²³<https://www.sciencenews.org/article/medical-advice-ai-chatbots-symptoms>

becomes standard, and physicians' independent diagnostic skills atrophy, the healthcare system becomes more, not less, fragile. Finally, HHS should support research on the informed consent gap: what patients actually understand about AI involvement in their care, how they evaluate AI safety claims, and what disclosure standards would support genuine informed consent. HHS should also fund research on community-partnered AI evaluation methodologies. The datathon model developed by MIT Critical Data provides structured opportunities for diverse stakeholders, including clinicians, data scientists, patients, and community members, to evaluate AI tools collaboratively. Current healthcare AI research overwhelmingly focuses on populations and conditions already well-served by existing infrastructure. HHS-funded research should explicitly prioritize AI applications addressing care gaps for rural, low-income, and historically underserved populations, and should require that evaluation protocols include the communities the AI is intended to serve.

XIV. Structural Recommendations

1. Establish federal minimum evidence standards for clinical AI.

Any AI system making autonomous clinical decisions should demonstrate safety and efficacy through independent, peer-reviewed clinical evidence before deployment. Company-funded preprints with equity-holding authors cannot constitute an adequate evidentiary basis. HHS should work with FDA to create a premarket review pathway for autonomous clinical AI proportionate to risk.¹²⁴

2. Close the CDS exemption loophole for functionally autonomous AI.

AI tools that functionally make clinical decisions, even those with nominal “clinician review” disclaimers, should not escape oversight through the 21st Century Cures Act exemption. HHS should clarify that the CDS exemption does not apply when AI recommendations are presented as the primary or default clinical action.¹²⁵

3. Require mandatory fairness testing.

Following Health Canada’s model, HHS should require clinical AI vendors to collect and report disaggregated performance data across demographic subgroups before and after deployment. The evidence that AI trained on non-representative data systematically harms minority patients is too extensive to ignore.¹²⁶

4. Establish a federal adverse event reporting system for clinical AI.

No coordinated federal mechanism exists for tracking adverse events from AI clinical tools outside FDA device regulation. HHS should establish a reporting system modeled on MedWatch for incidents involving non-device clinical AI.¹²⁷

5. Issue HIPAA guidance for autonomous AI systems.

HHS should clarify that autonomous AI systems require explicit patient authorization before accessing prescription data, separate from the general TPO exception. The current framework, in

¹²⁴<https://pmc.ncbi.nlm.nih.gov/articles/PMC12140231/>

¹²⁵<https://bipartisanpolicy.org/issue-brief/fda-oversight-understanding-the-regulation-of-health-ai-tools/>

¹²⁶<https://www.canada.ca/en/health-canada/services/drugs-health-products/medical-devices/application-information/guidance-documents/pre-market-guidance-machine-learning-enabled-medical-devices.html>

¹²⁷<https://home.ecri.org/blogs/ecri-news/misuse-of-ai-chatbots-tops-annual-list-of-health-technology-hazards>

which patient data flows from pharmacy to Surescripts to commercial AI startup without notification, is inconsistent with HIPAA's privacy protections.¹²⁸

6. Require transparency in state AI sandbox programs involving clinical care.

Safety data, adverse event reports, and clinical evaluation methodologies from state sandbox programs must be publicly accessible. Business confidentiality cannot override public health transparency. HHS should condition federal cooperation with state sandboxes on minimum transparency standards.¹²⁹ Additionally, HHS should require a standardized "Model Facts Label" for all clinical AI tools, analogous to a nutrition label, that discloses in plain language: the training data composition and known gaps, validated performance metrics across demographic subgroups including rural and lower socioeconomic populations, known limitations and failure modes, update history and version control, and the developer's conflict of interest disclosures. This label should be accessible to both clinicians and patients at the point of care.¹³⁰

7. Develop a liability framework for autonomous clinical AI.

HHS should convene stakeholders to develop clear liability apportionment for autonomous AI clinical decisions, considering strict liability models, common enterprise liability, and compensation funds. The current vacuum, where neither malpractice nor product liability clearly applies, leaves patients without recourse.¹³¹

8. Require disclosure of funding sources and conflicts of interest.

Clinical AI companies citing studies to support clinical claims should be required to disclose funding sources, author equity interests, and peer-review status. HHS should examine whether VC firms providing both investment and regulatory strategy create conflicts that compromise regulatory integrity.

9. Establish a federated AI Sentinel network for post-market surveillance.

Modeled on FDA's existing Sentinel Initiative for drug safety, HHS should fund a federated AI Sentinel network that collects standardized performance and safety signals from participating health systems. This network should monitor for performance drift, subgroup degradation, and emerging safety signals across deployed clinical AI tools, and publish de-identified safety bulletins and recall-like advisories when models underperform or cause harm. The current

¹²⁸<https://www.law.nyu.edu/sites/default/files/Comment%20of%20NYU%20TLP%20Clinic%20on%20Proposed%20HIPAA%20Rule.pdf>

¹²⁹<https://www.nature.com/articles/s41591-024-03304-x>

¹³⁰GRAMA Appeal Response, Utah Dept. of Commerce, Feb. 18, 2026 (on file with authors).

¹³¹<https://www.ncbi.nlm.nih.gov/books/NBK613216/>; https://repository.law.umich.edu/cgi/viewcontent.cgi?article=1569&context=book_chapters

post-market surveillance infrastructure is inadequate: the FDA's MAUDE system captures fewer than 0.5% of adverse events, and Predetermined Change Control Plans allow AI modifications without fresh clearance. A dedicated Sentinel network would provide the continuous monitoring infrastructure that clinical AI requires.¹³²

¹³²<https://www.fda.gov/safety/fdas-sentinel-initiative>

XV. Conclusion

This RFI asks how HHS can “accelerate the adoption and use of AI as part of clinical care.” We have documented in this comment why acceleration without evidence standards does not constitute responsible innovation. The evidence base for the most prominent clinical AI systems is company-funded and unreviewed. Documented patient harms range from eating disorders to teen suicides to systematic racial bias. The regulatory framework has gaps that no federal agency has claimed. The liability system cannot tell patients who is responsible when AI harms them. And the financial incentives driving deployment are structurally misaligned with patient safety.¹³³

The pattern documented in this comment, in which companies deploy AI with minimal validated evidence, navigating regulation through political strategy rather than clinical proof, access patient data through frameworks never designed for autonomous AI, and generate claims that major media amplify without scrutiny, is not confined to one company or state. It reflects systemic pressures that will only intensify.

Every major medical society opposes autonomous AI prescribing without rigorous safeguards. Every comparable international jurisdiction imposes more protective requirements. The documented record of AI health tool harms, including eating disorder patients receiving dieting advice, teenagers dying by suicide, and algorithms that miss 67% of sepsis cases, demonstrates the technology is not ready for autonomous clinical deployment without robust protections.

The Department has both the authority and the responsibility to establish the federal floor. Doing so will not slow innovation. It will build the public trust that is the precondition for AI’s sustainable integration into American healthcare.

We appreciate the opportunity to comment on this important matter. We are available to provide additional information, testimony, or briefing at the Department’s request.

Respectfully submitted,

MIT Critical Data

<https://criticaldata.mit.edu/>

¹³³<https://www.medrxiv.org/content/10.1101/2025.07.14.25331406v1>