radiology partners

February 23, 2026
Department of Health and Human Services (HHS)
Office of the Deputy Secretary and Assistant Secretary for Technology Policy (ASTP)
Office of the National Coordinator for Health Information Technology (ONC)
Attention: HHS Health Sector AI RFI

Re: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

Radiology Partners appreciates the opportunity to submit comments to the request for information on "Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care" (Docket No. HHS-ONC-2026-0001-0001) published by the Office of the Deputy Secretary and Assistant Secretary for Technology Policy (ASTP) and Office of the National Coordinator for Health Information Technology (ONC) (collectively, ASTP/ONC), Department of Health and Human Services (HHS).

Radiology Partners, through its affiliated practices, is a leading technology-enabled radiology practice in the United States, serving more than 3,400 healthcare facilities with high quality radiology and technology, including multiple artificial intelligence (AI) solutions. As a physician-led and physician-owned practice, our mission is to transform radiology by innovating across clinical value, technology, service and economics, while elevating the role of radiology and radiologists in healthcare. Radiology Partners, through its affiliated practices, provides consistent, differentiated care to patients, while delivering enhanced value to the hospitals, clinics, imaging centers and the referring physicians we serve.

Radiology is at the cutting edge of AI technology, serving as one of the earliest medical specialties to integrate AI into clinical practice. For nearly a decade, radiologists have used AI tools to assist with image interpretation and streamline workflows. This early and ongoing integration has positioned radiology as an example for other fields seeking to integrate AI effectively. Over the past few years, Radiology Partners has led the largest clinical deployment of AI in radiology and healthcare more broadly. With this depth of experience, we have continually refined our processes for developing, validating, deploying, and monitoring AI models and welcome the opportunity to share these insights with HHS.

We address HHS' requests for information on AI regulation, reimbursement, and research & development and its specific questions through the lens and experience of a radiology practice that has deployed AI models at scale. While radiology-specific examples are used, the principles and processes described apply broadly across clinical specialties.

<u>Specific Questions:</u>

Question 1: What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

Regulatory Barriers: One of the most significant barriers to AI innovation and adoption is the misalignment between current regulatory frameworks and how modern AI systems function in real-world clinical care, particularly the reliance on finding-level, pre-market validation requirements that do not scale to foundation or generative models designed to interpret entire imaging studies across over a thousand possible findings. The current regulatory system disproportionately rewards AI discovery over delivery and does not support the infrastructure required for safe and scaled AI deployment in the hospital ecosystem. In addition, the regulatory framework relies heavily on pre-market clinical validation methods conducted in controlled settings that often fail to reflect real-world performance, usability, and workflow integration in the clinical environments where tools are deployed. While these studies are intended to ensure safety and efficacy before clinical use, they often do not account for the real-world variability in patient populations, imaging protocols, exam types, clinical workflows, and physician interpretation. As a result, the current approach can slow innovation and inflate costs without proportionally improving safety because it relies on controlled test performance rather than longitudinal performance monitoring in live clinical environments.

Infrastructure Gap: Many hospitals and clinical settings lack the technical, operational, and human capital systems required to implement AI and other advanced technologies. They need:

- Modern data infrastructure to handle imaging and clinical data at scale, both of which are unstructured;
- Interoperability across hospital technology and data platforms including electronic health record (EHR) and imaging systems (e.g., Picture Archiving and Communication Systems (PACS), Radiology Information System (RIS));
- Dedicated expert physician teams for AI validation, continuous monitoring, education, and change management; and
- Formal partnerships with vendors to ensure transparency and clinical usability of the AI models at the point of care.

Currently, each hospital must independently build this infrastructure, with no funding and few guidelines. Without a clear reimbursement model or implementation roadmap, most health systems cannot afford to engage and deploy AI solutions, even when the innovation has clear, clinical promise. Many AI tools improve efficiency, quality, and patient safety but do not align with payment models, leaving provider groups and hospital systems to absorb implementation, validation, and monitoring costs without corresponding financial support. This misalignment discourages investment in infrastructure and disproportionately limits adoption of AI tools to well-resourced systems.

Data Availability and Standardization Challenges: Radiology data is often highly heterogeneous, with wide variation across healthcare systems, protocols, PACS configurations, and reporting systems. These differences create meaningful barriers during AI development, deployment, integration, and monitoring. In addition, there are critical gaps in data availability, particularly robust pediatric imaging datasets paired with accurate, clinically detailed radiology reports, limiting the ability to develop and validate robust AI models in these populations. While large language models (LLMs) can help extract and structure information from free-text reports, they are not sufficiently consistent, transparent, or reliable to replace standardized data capture and governance in high-stakes clinical environments. This work remains resource-intensive, particularly when supporting thousands of facilities with heterogeneous and evolving data environments. Without broader adoption of data standards, benchmarking frameworks, and interoperability requirements, data variability and data scarcity will remain major barriers to innovation and generalizable AI performance.

AI Model Transparency: A critical barrier to AI adoption is insufficient user-facing transparency and traceability of the AI model at the point of care. When physicians do not have access to essential contextual information, such as intended use, known limitations, and any clinically meaningful impacts, physicians may struggle to appropriately calibrate trust in the AI's output. This lack of visibility increases the risk of human-AI interaction errors, including automation bias (accepting AI false positive results) and complacency (accepting AI false negative results), and can undermine physician confidence in the tool. These transparency gaps make it difficult for physicians to distinguish between true model output, missing or pending output, and inferred content, ultimately slowing adoption and limiting safe, reliable use of AI in routine clinical workflows. HHS could encourage public trust by mandating a "model card" system whereby a tool could provide a clear summary of its performance across granular demographics (e.g., age, gender, socioeconomic status). This would allow clinicians to determine quickly if a model is trustworthy for use in local patient populations and would create an equitable solution that allows a range of institutions to decide which AI solutions best fit their needs and create trusted toolkits.

Human Factors and Physician Trust Barriers: AI adoption depends on physician trust and appropriate trust calibration. Distrust can be common early in deployment, particularly with novel vendors, unfamiliar workflows, or new model types, and may lead radiologists to dismiss accurate clinically impactful AI findings. Over time, as familiarity grows, the risk can shift toward over-reliance on the AI tool, including automation bias and complacency. Supporting safe and effective use requires more than technical model performance; it requires implementation practices that help clinicians understand how the tool fits into their clinical workflow, what specific value it provides, and where it may underperform. This calibration requires initial training, transparent communication of limitations and failure modes, and ongoing monitoring and feedback. Importantly, governance frameworks should support monitoring not only of the accuracy of the AI model output, but also clinician interaction patterns (e.g., override rates, reliance trends, and recurrent error types), as these behaviors directly influence patient safety. Trust can be lost quickly after a single, unexpected failure, particularly when failure modes are not transparent, anticipated, or detectable at the point of care.

Question 2: What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.

Regulatory Policy Changes:

1. ***Enable Life-Cycle Based Oversight for Clinical AI:*** Current regulatory pathways emphasize pre-market clinical validation in controlled settings that do not reflect real-world variability and are not scalable to modern foundation and generative models. HHS should collaborate with the FDA to establish a lifecycle-based oversight framework that explicitly accommodates iterative model updates and continuous learning systems. This could take the form of:
    - A conditional authorization pathway that permits deployment with a defined post-market surveillance plan (analogous to EU's *Software as a Medical Device (SaMD) lifecycle approach*).
    - Requirements for real-world performance reporting (case mix, population subgroups, workflow settings) as part of periodic reauthorization.
    - Publication of performance drift and safety metrics as part of post-market reports.

Relevant policy anchor:

- 21 CFR Part 807 (Establishment Registration and Device Listing) and 21 CFR Part 820 (Quality System Regulation) could be extended with lifecycle surveillance expectations appropriate for clinical AI.
- FDA's "Total Product Lifecycle" program for certain devices already conceptually supports structured post-market activities; this should be expanded to cover AI/ML systems.

There is currently no FDA-defined, standardized postmarket monitoring framework for radiology AI, which creates challenges for both vendors and providers. A successful approach must avoid creating unnecessary documentation burdens or encouraging unvalidated, adhoc reporting. At the same time, responsibility for monitoring should not rest solely with vendors. Because radiology AI tools are used under the supervision of U.S. board certified radiologists, real world oversight naturally occurs as providers adjust how often they use a tool when the tool's performance does not meet expectations. Accordingly, any federal framework should focus on structured, validated mechanisms for postmarket monitoring that balance accountability with practicality and ensure that feedback is clinically meaningful. This approach would support safety while keeping monitoring feasible and proportional to real world clinical workflows.

2. ***Mandate Transparency and Explainability Standards at the Point of Care:*** To improve clinician understanding and mitigate human–AI interaction errors (automation bias, complacency, under-reliance), the FDA should require user-facing transparency at the point of care. AI outputs should not be limited to raw predictions; they must include contextual information such as:
   - Model version and release notes.
   - Confidence scores.
   - Images and data elements used in inference (which series, which slices for the current and, if applicable, prior exam).
   - Explicit identification of excluded studies and reasons (e.g., outside cleared indications).
   - Status information (received, in process, complete, rejected).

These elements should be standardized and interoperable using established frameworks such as Integrating the Healthcare Enterprise (IHE) AI profiles.

Relevant policy anchor:

- 21 CFR 803 (Medical Device Reporting) could be supplemented with requirements for contextualized outputs to improve traceability and post-market auditability.

3. ***Harmonize Cross-Agency Data Standards to Reduce Redundancy:*** From the perspective of a radiology group responsible for deploying AI across diverse clinical settings, the lack of consistent definitions and data standards across CMS, FDA, and ONC leads to duplicative burdens for providers and limits safe AI implementation. HHS should initiate a coordinated cross-agency effort to align terminologies, metadata requirements, and semantic definitions for clinical data and imaging standards (e.g., RadLex, LOINC, SNOMED, DICOM tags). This effort would reduce variability and support AI generalizability.

Relevant policy anchor:

- 45 CFR Part 170 (Health IT) and ONC Health IT Certification Program criteria should explicitly incorporate semantic standard requirements for AI readiness.

4. ***Modernize HIPAA for AI-Relevant Use Cases:*** Current HIPAA de-identification requirements obstruct multi-institutional aggregation of rich datasets for training, validation, and ongoing monitoring. HHS should issue updated guidance or rulemaking clarifying the following:
   - Safe, standardized de-identification for AI research and validation.
   - Conditions for limited datasets that include necessary metadata elements (protocols, demographics) while protecting privacy.

Relevant policy anchor:

- 45 CFR Part 164 Subpart E (Privacy of Individually Identifiable Health Information), particularly §§164.514(b)–(c), should be revisited with AI research requirements in mind.

Payment Policy:
Current payment models are structurally misaligned with AI-enabled care. Today's Medicare Physician Fee Schedule (PFS) evaluates and reimburses discrete procedural tasks. On the other hand, AI models often require radiologists to review, validate, and adjudicate AI-generated insights, increasing effort and cognitive load, creating value that the current PFS cannot recognize or reimburse appropriately.

The current budget neutrality requirement within the PFS is a major structural barrier to incorporating AI-enabled cognitive work into the Medicare payment system. Because new services must be offset by cuts elsewhere, introducing AI-enabled cognitive work into the PFS would force reductions in unrelated physician services. For AI to deliver meaningful clinical benefit, HHS must consider alternative payment approaches outside the PFS and outside traditional payment mechanisms.

5. ***Support Infrastructure and Adoption Through Incentive Programs**:* AI deployment requires robust infrastructure (secure data pipelines, monitoring processes, clinician training) that is not currently reimbursable. HHS should create a voluntary, incentive program for AI adoption that funds:
   - Data infrastructure upgrades (scalable de-identification, secure inference pipelines).
   - Local validation and governance workflows.
   - Continuous monitoring systems that capture real-world performance metrics.
   - AI training.

This burden is especially heavy when hospitals are expected to fund these changes based solely on potential, indirect return on investment (ROI), such as reduced length of stay. Given the foundational limitations of the PFS, federal investment should reflect the reality that advanced AI tools reshape clinical workflows, clinical decision-making, and diagnostic quality by rewarding value-based outcomes.

A federal incentive program modeled could accelerate adoption, fund necessary infrastructure, and establish baseline best practices for efficient, safe and effective AI deployment. By supporting foundational infrastructure, such a program can mitigate reliance on direct reimbursement for individual AI models, bypassing the complexities of the existing healthcare payment system while supporting AI adoption, guided by best practices. It would also reward hospitals for adopting best practices in AI deployment, including building the necessary infrastructure and processes for safe and scalable implementation, optimization of physician-AI collaboration, and maintenance of model accuracy over time.

Possible policy anchor:
- CMS Innovation Center (CMMI) authority to test new payment and delivery models could be used to pilot AI infrastructure support programs.

Programmatic Design:
6. ***Integrate Explainability and Workflow Design into Performance Goals**:* For AI to deliver clinical value, tools must integrate seamlessly into radiologist workflows. Programmatic support should encourage:
   - Usability testing and human factors engineering as part of AI validation.
   - User interfaces that present relevant contextual information without overload (e.g., confidence flags, region highlights).

- Explainability outputs tailored to clinical needs (categorical likelihood cues rather than opaque numeric scores).

7. ***Mandate Ongoing User Training, Education, and Feedback:*** Effective AI use goes beyond initial buttonology. HHS should support programs for:
    - Continuous clinician education tied to model updates.
    - Feedback loops between clinicians and developers.
    - Monitoring of clinician–AI interaction patterns (override rates, drift in reliance).
    - Conditioned funding tied to evidence of safe and appropriate adoption.

Clinical education focuses on the AI model's capabilities, limitations, and appropriate use within the radiologist's workflow. Training is most effective when it is continuous, with refreshers provided when new model versions are deployed or when monitoring metrics indicate that a radiologist may be at increased risk for distrust, automation bias or complacency. We have observed cases in which hospitals implemented AI tools, but radiologists were unaware of their availability or were uncertain how to use them effectively. Ongoing clinical education ensures clinicians understand both the value of the tool and its boundaries. Group training, videos, and in-app education are effective mechanisms for delivering continuous training at scale, but they cannot yet fully replace initial, individualized instruction for more complex tools.

8. ***Standardize Communication and Feedback Mechanisms:*** Program design should support multiple communication modalities tied to clinician workflows (in-app, dashboards, alerts) to reinforce updates and performance insights, which is a best practice for safe adoption.

Why These Changes Matter: Taken together, these regulatory, payment, and programmatic reforms would:

- Align oversight with real-world clinical environments.
- Reduce administrative burden through harmonized standards.
- Incentivize hospital investment in safe, scalable AI infrastructure.
- Improve clinician trust and appropriate use through transparency and training.
- Enable continuous monitoring and longitudinal evaluation of real-world performance.

Question 3: For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (*e.g.,* relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?

Non-medical device AI tools, including generative AI used for documentation and workflow support, may create novel governance challenges because accountability, liability, and oversight expectations are not well defined. Current frameworks place traditional malpractice liability at odds with product-liability pathways, creating uncertainty about how responsibility should be allocated when AI outputs influence clinical decision-making through triage or administrative functions. These gaps are compounded by inconsistent vendor indemnification terms, limited transparency into model updates and data handling, and insufficient auditability (e.g., versioning, traceability, and logging of inputs used to generate outputs). Privacy and security concerns are also heightened when protected health information is transmitted to external AI services, and existing privacy frameworks were developed before today's large-scale analytics and foundation models.

HHS can help by issuing fit-for-purpose guidance and baseline requirements that enable safe implementation without imposing burdensome pre-market documentation or slowing clinical deployment. As part of this, HHS should work with the Office of the General Counsel to create a safe harbor for algorithm transparency, such as providing federal guardrails that clarify responsibility for AI-influenced decisions while ensuring clinicians remain the final authority and AI serves as a supportive tool. HHS should also modernize privacy and security expectations by developing interoperability and data-handling standards that move beyond basic encryption to include federated learning approaches and national de-identification landmarks. This “gold standard” would strengthen data stewardship while reducing the cost of repetitive, vendor-specific security audits.

Finally, HHS should consider establishing a unified AI risk framework grounded in “safety-by-design” principles. Such a framework would differentiate requirements for high-risk clinical AI versus lower-risk administrative tools, reducing unnecessary compliance burden while ensuring appropriate safeguards. By clarifying liability, strengthening privacy protections, and aligning oversight with actual risk, HHS can support scalable, clinically responsible adoption of non-medical device AI while maintaining clinician-led decision-making and patient safety, allowing organizations to focus on clinical outcomes instead of administrative compliance.

Question 4: For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (*e.g.,* contracts, grants, cooperative agreements, and/or prize competitions)?

For non-medical device AI tools used in clinical care, such as generative AI used for documentation and workflow support, the most promising evaluation methods focus on real world usability, reliability, human factors behavior, and longitudinal performance, rather than traditional medical device validation metrics. Models deployed in live environments encounter real-world variability due to differing hospital settings and scanner types which do not match the testing environment. To harden models against real-world variability and diagnostic drift, HHS should consider prioritizing standardized robustness training, including variations in documentation style, clinical vocabulary, exam types, and workflow interruptions. HHS should encourage joint governance model centered on the Chief Medical Information Officer (CMIO) and Chief Quality Officer (CQO)/Quality & Safety committees. Evaluation and monitoring should extend beyond simple accuracy and sensitivity metrics to focus on overall clinical utility in real-world scenarios.

Measuring human-centered safety metrics help mitigate diagnostic drift as the model experiences changing clinical protocols or gets exposed to varying patient populations. Safety metrics, such as override rates and alert fatigue levels, can be used to assess whether clinicians are using the tool as a true decision support tool or if the tool is a distraction adding to the overall cognitive load, which could lead to “near miss” events and patient safety concerns. Measuring real-time discordance rates between AI outputs and human assessment can proactively identify performance gaps and biases allowing quality and safety teams to proactively mitigate against any effects on patient care ensuring that “Human plus AI” model remains a robust safeguard. HHS can accelerate this by funding standardized safety dashboards through grants and cooperative agreements, enabling CMIO/CQO leaders to monitor performance across sites using actionable indicators such as alert-to-action ratios (how often AI outputs actually change management). Support for these standardized approaches should occur through cooperative agreements that provide the AI industry with standardized approaches that center on providing positive impact on the clinician's decision making and not on technical speed of the enterprise.

Question 5: How can HHS best support private sector activities (*e.g.,* accreditation, certification, industry-driven testing, and credentialing) to promote innovative and effective AI use in clinical care?

HHS can accelerate safe and effective AI adoption by supporting private-sector accreditation, certification, and credentialing programs that address the gap between FDA authorization and real-world clinical use. While FDA review evaluates whether an AI tool is safe and effective for its intended use, real-world performance and patient safety are strongly influenced by implementation quality, workflow integration, clinician training, and ongoing monitoring. These factors are not consistently addressed today and vary widely across healthcare systems and vendors. To address this, HHS can support the private sector by moving away from static certification models to continuous clinical accreditation. One-time approval does not account for inevitable diagnostic drift of the model as it deploys in real-world scenarios and could negatively affect disparate patient populations without safeguards in place. Therefore, we recommend HHS incentivize models that demonstrate and prioritize real-time ongoing safety monitoring. By supporting such platforms, HHS can ensure that AI tools remain continuously safe and reliable throughout the entire lifecycle of the product.

A high-impact opportunity is defining and supporting credentialing for qualified end users of healthcare AI. For imaging AI, safe use depends on clinicians who understand the model's intended use, limitations, common failure modes, and appropriate reliance behaviors. Existing professional society guidance, including the American College of Radiology's (ACR) AI policy, emphasizes the importance of a specialist human in the loop and appropriate user training. Beyond end-user training, innovation is also constrained by the lack of standardized, high-quality reliable datasets for testing AI tools prior to deployment. To address this barrier, HHS could consider cooperative agreements to fund secure, de-identified data banks where responsible AI developers can test model performance against real-world data from multiple institutions. Federal support for these shared testing environments would reduce barriers to innovation, maintain a high bar for performance equity, and shift industry practices toward a "safety-first" environment where evaluation centers on real patient outcomes rather than idealized laboratory datasets. Supporting the development of standardized competency expectations for AI users and representative testing frameworks would strengthen accountability and ensure AI tools are safe and effective in clinical settings, all in alignment with CMS quality frameworks and centered on real patient outcomes.

In parallel, HHS could support organizational-level accreditation models for AI governance maturity. This would recognize health systems that implement core best practices, such as local validation, transparency and auditability, post-deployment monitoring, incident reporting, and ongoing education. HHS support could include grants, cooperative agreements, and public-private partnerships to develop standardized evaluation frameworks, shared monitoring approaches, and scalable training resources.

In addition, HHS could encourage development of new professional credentialing standards focusing on Chief Medical Information Officers (CMIOs) and Quality/Safety committees in ensuring the clinical effectiveness of AI models. Proper use of AI tools is a true clinical skill, not just technical, and credentialing should therefore include assessment of automation bias and methods to manage "human plus AI" workflows. Specifically, HHS could consider awarding grants to professional societies to develop explainability standards and training modules that support responsible AI stewardship. This would help shift the culture away from IT-centric deployment toward clinician-led human-factors engineering, improving adoption, trust, and safety. Such an approach could elevate, rather than replace, the professional abilities of board-certified clinicians by ensuring AI serves as a supportive tool within well-governed clinical workflows.

Question 6: Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

In our experience, the AI tools that have most consistently met or exceeded performance and cost expectations are foundation-model workflow tools that help manage today's exponentially higher clinical intensity and cognitive load due to high resolution multisequence datasets. In radiology, this includes generative AI systems that support end-to-end reporting, not only narrow single-finding algorithms, and synthesize thousands of imaging data points into coherent, actionable insights. By reducing cognitive burden and structuring information at the point of interpretation, these tools enable greater diagnostic precision, consistency across variable studies, and clearer communication of clinical reasoning.

Two foundation-model tools have been particularly impactful. Mosaic Reporting, an ambient reporting tool built on large language models, enables radiologists to create structured reports with fewer words, dictate in any order, and use natural language commands and embedded macros within templates. Mosaic Drafting, a home-grown vision-language model (VLM), generates a full draft report for the entire study (including comparison to priors), functioning like an "AI resident" that shifts the radiologist from report creation to oversight and validation.

AI tools most commonly fall short when implementation and human factors are underestimated, including poor workflow integration, limited transparency into model behavior, and insufficient training and monitoring to prevent under-use or over-reliance. The most promising next-generation tools will combine foundation-model capabilities with robust post-deployment monitoring, transparency, and governance frameworks to ensure safe and scalable real-world performance.

HHS should consider prioritizing the following:

- Emphasize AI's role in managing data complexity to ensure diagnostic precision and move away from pure time savings metrics.
- Recognize prospective drift detection and mitigation to avoid downstream costs from missed or delayed diagnoses
- Support payment and research models that recognize the cognitive effort required by an expert to validate AI generated analysis and insight and that value diagnostic certainty as a measure of quality.

Question 7: Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

Within health care organizations, the roles most critical to successful AI adoption are those that bridge clinical leadership and technology oversight. Historically, AI implementation has been driven primarily by IT departments, which are important backbones to organizations but are not always positioned to evaluate clinical utility, workflow fit, or patient-centered outcomes. To accelerate responsible adoption, organizations should shift decision-making toward combined clinical leadership models, where roles, such as the Chief Medical Information Officer (CMIO), Chief Quality Officer (CQO), and relevant clinical department leaders jointly oversee AI evaluation, deployment, and governance. This governance framework ensures that AI investments are aligned with patient-centered outcomes and professional workforce needs rather than driven by hardware. Empowering clinician-led governance also strengthens trust and supports thoughtful adoption grounded in real-world care, not technical capability alone.

Radiologists today face significant cognitive overload due to fragmented workflows (i.e., the need to switch between multiple AI interfaces layered on top of the primary diagnostic workstation). Such fragmentation can erode efficiency, increase error risk, and stall adoption. To reduce these administrative hurdles, HHS should consider promoting deep workflow integration models that embed AI tools directly into existing clinical systems. Because current vendor ecosystems rely on non-standardized Application Programming Interfaces (APIs), organizations often struggle with interoperability barriers that prevent seamless integration of third-party AI outputs. HHS could support policies that standardize universal API requirements to reduce interoperability friction and penalize vendors who restrict the incorporation of third-party AI insights into diagnostic viewers and cause blockage of information transfer.

A final administrative hurdle comes from the cumulative burden of contracting, security reviews, and legal negotiations required each time a hospital adopts a new AI tool. These repetitive processes can delay deployment of life-saving AI capabilities despite clear evidence of patient benefit. To streamline adoption while maintaining strong safeguards, HHS could consider a platform model where hospitals gain access to a vetted “AI gateway” under a master agreement. From this gateway, organizations could select among multiple AI tools without initiating a full contracting and security cycle for each product. This approach would reduce paperwork, improve speed-to-implementation, and ensure that both large academic centers and smaller rural hospitals can equitably access advanced AI tools without disproportionate administrative burden.

Question 8: Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

HHS should recognize interoperability as a mandatory safety guardrail for clinical AI. In the current framework, AI outputs are often trapped in application or vendor-specific silos, which limits auditability, hinders cross-site learning, and weakens longitudinal safety monitoring. To accelerate reliable, equitable adoption, policy should shift from episodic data exchange between healthcare entities to semantic liquidity where AI outputs and supporting context are encoded in shared, computable formats that flow with the patient and remain actionable across systems, while remaining fully compliant with HIPAA privacy and security requirements. HHS should also consider the use of Fast Healthcare Interoperability Resources (FHIR) native AI results, which would make outputs versioned, auditable, and actionable within EHR workflows, strengthening traceability and supporting consistent physician use. Linking imaging metadata with longitudinal outcomes can prevent diagnostic drift.

HHS can also work with the FDA to allow “safe harbors” for organizations to aggregate de-identified data, which would allow for the development of national benchmarks for AI performance and safety reflecting real-world diversity without increasing regulatory burden.

There is an opportunity for HHS and the FDA to support a framework for continuous performance monitoring where developers can safely test their models against real-time institutional data. This approach would give developers visibility into how models perform across diverse patient populations, imaging equipment, and care settings, ensuring that health equity is a focus on initial development. Overall, HHS can complement its current focus on interoperability by emphasizing clinical quality and safety that promotes the development of reliable, equitable tools.

Question 10: Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?

a) **Are there published findings about the impact of adopted AI tools and their use clinical care?**
b) **How does the literature approach the costs, benefits, and transfers of using AI as part of clinical care?**

HHS should prioritize AI research that reflects real-world clinical environments rather than static, historical test conditions. Traditional validation using retrospective datasets does not capture the variability introduced by differences in patient populations, imaging protocols, scanners, workflows, and clinical context. To accelerate safe and effective adoption, HHS should support prospective, multi-center clinical trials that evaluate AI tools in live workflows, including how clinicians interact with them in routine care. This research should emphasize human-in-the-loop integration, evaluating how AI maximizes diagnostic accuracy and improves turnaround times when used as a dynamic member of the care team. This research would allow for HHS to establish long-term evidence for reimbursement and increase trust for patients and providers

In addition, HHS should emphasize research into human factors engineering, including how User Interface (UI) design influences clinician's critical thinking; this would mitigate long-term effects of automation bias. An example of this research includes measuring the optimal amount of data that should be available to the clinician without causing "alert fatigue" and information overload. This research would ensure that the AI remains a clinical decision support tool and not a replacement for clinician judgment.

We also recommend that HHS develop a national continuous quality performance monitoring platform to prevent diagnostic drift and algorithmic bias across wide ranges of demographics and imaging hardware. HHS can also consider developing a national benchmarking clearinghouse to help ensure equitable AI performance, regardless of local technology or demographics. Overall, a coordinated federal research agenda, spanning clinical trials, human-factors studies, performance-drift monitoring, and national benchmarking, would establish the evidence base needed to guide reimbursement, strengthen clinician and patient trust, and ensure AI is deployed safely, equitably, and effectively across the healthcare system

Conclusion

Radiology Partners appreciates the opportunity to provide comments in response to HHS' request. As HHS advances its "OneHHS" AI Strategy and seeks experience-based input to guide regulation, reimbursement, and research, radiology's long-standing use of AI offers practical insights to support this effort. We respectfully encourage HHS to continue engaging with the radiology community as HHS advances a forward-leaning, industry-supportive, and secure approach to accelerate the adoption and use of AI. We look forward to continued partnership as HHS refines federal policy to advance responsible AI adoption across clinical care.

Respectfully Submitted,

Dr. Nina Kottler, MD, MS, FSIIM, FAIM
Chief Medical AI Officer, Mosaic Clinical Technologies @Radiology Partners

Justine Jamero
Associate Director, Health Policy & Advocacy