

Ratnalekha Viswanadham
Cupertino, CA
February 23, 2026

U.S. Department of Health and Human Services
200 Independence Avenue, SW
Washington, DC 20201

Re: Request for Information – Accelerating the Adoption and Use of Artificial Intelligence as Part of Health Care Delivery

To Whom It May Concern:

I appreciate the opportunity to comment on how HHS can accelerate the adoption and use of artificial intelligence (AI) in health care delivery.¹ My comments primarily address Questions 1, 2, 8, 9, and 10 in the Request for Information, with particular emphasis on barriers to adoption (Q1), regulatory and program design changes (Q2), data and measurement infrastructure (Q8), patient- and caregiver-relevant challenges (Q9), and priority research and evaluation needs (Q10).

I write as a neuroscientist, a health services and digital health researcher who has worked in health care settings, and as a former detailee under an Intergovernmental Personnel Agreement for federal evaluation and evidence-building efforts with the U.S. General Services Administration's Office of Evaluation Sciences (OES).² In my academic research, I focus on how electronic health record (EHR) data and EHR use metadata can be used to understand health care delivery, and on how digital tools such as telemedicine have diffused across different organizations and communities. My experience at OES has also emphasized the importance of embedding simple, pragmatic evaluation designs into real-world programs so that agencies can learn what works while they implement. I also bring nearly two decades of experience as an educator, having studied learning and cognitive load and developed strategies to help students learn material in challenging circumstances (e.g., learning disabilities, English as a second language, high academic pressure, and under-resourced settings) globally.

Below, I first offer a concise set of policy-oriented recommendations, followed by a more detailed discussion. In this response, I use "AI" to denote data-driven learning systems (such as statistical and machine-learning models, large language models, generative AI, agentic AI, and other adaptive decision-support tools). This is because it is the learning element, rather than general rule-based automation, that raises unique questions regarding evaluation, drift, and oversight.

Summary of recommendations (Q1, Q2, Q8, Q9, Q10)

1. Require continuous evaluation for any federally supported AI implementation.
Make evaluation a non-negotiable component of AI deployment, aligned with the Evidence Act's intent, with clear expectations for what is measured, how often, and how results are used to adjust or even pause implementation (Q1, Q2, Q8, Q10).
2. Build an evaluation framework that uses EHR and EHR-use metadata.
Use existing EHR metadata (e.g., audit logs, documentation patterns, workflow traces) to measure how AI affects workflow, team dynamics, and the cognitive environment, rather than just whether the AI is on (Q1, Q2, Q8, Q10).³

¹ <https://www.federalregister.gov/documents/2025/12/23/2025-23641/request-for-information-accelerating-the-adoption-and-use-of-artificial-intelligence-as-part-of>

² <https://oes.gsa.gov/>

³ <https://www.jmir.org/2025/1/e64721>

3. Tie funding and incentives to demonstrated value, not just use
Link grants, pilots, and payment adjustments to improvements in a small set of meaningful, feasible, and viable measures (for example, burden reduction, timeliness, care coordination), rather than to raw AI adoption rates (Q1, Q2, Q10).⁴
4. Focus on “automation of the obvious, flagging of the not-so-obvious.”
Encourage AI that automates routine, well-understood tasks and surfaces non-obvious risks, while preserving human expertise and judgment for complex and atypical cases (Q1, Q2, Q9, Q10).
5. Guard against automation complacency and overload.
Require designs and training that keep humans actively engaged, monitor for over-reliance and new types of error, and ensure clinicians and staff can operate safely even when AI systems fail or are unavailable (Q1, Q2, Q9, Q10).⁵
6. Support adoption in resource-constrained and non-tech-hub settings
Pair AI incentives with infrastructure and technical assistance for rural, small, safety-net, and other resource-limited organizations, so AI does not deepen the divide between large, well-resourced systems and everyone else (Q1, Q2, Q8, Q9).
7. Close the “model-to-practice” gap for federally funded AI.
Encourage NIH and other HHS funders to support not only model development but also implementation, evaluation, and operations partnerships, so that fewer AI models end at publication and more are translated into tools that improve care in practice (Q1, Q2, Q8, Q10).⁶

1. Continuous evaluation as a condition of federal support (Q1, Q2, Q8, Q10)

This section responds to Questions 1, 2, 8, and 10 by recommending that evaluation and learning be built into AI deployment from the start, using existing data streams where possible.

AI in health care is being introduced into complex systems where changes in one part of the workflow can have unintended effects elsewhere. A one-time pre-deployment assessment is not enough. To align with the spirit of the Evidence Act, HHS should make continuous evaluation a core requirement for any AI tool receiving federal funding or deployed under a federally supported program.

Key elements could include:

- Clear evaluation plans from the outset: Any proposal involving AI should include a brief but concrete evaluation plan outlining what will be measured, why those metrics matter, how data will be collected (e.g., from EHR metadata), and how often results will be reviewed. This can be scaled to project size, but it should never be optional.
- Use of routine data streams: Where possible, evaluation should use data already generated during care, particularly EHR and EHR-use metadata. This minimizes burden and makes it realistic to continue evaluation over time.
- Feedback loops, not just reporting: Programs should require that evaluation results be fed back into local decision-making. For example, if an AI tool appears to increase documentation burden or error rates, there should be a process to adjust the configuration, provide additional training, or suspend use.
- Guardrails for harms: HHS can set expectations that continuous evaluation includes monitoring for negative trends (for example, increased after-hours work, alert fatigue, or widening gaps between organizations) and that there are clear criteria for pausing or revising deployments when harms are detected.

⁴ <https://www.jonesday.com/en/insights/2026/02/hhs-signals-policy-direction-to-accelerate-adoption-of-ai-in-clinical-care>

⁵ <https://ojs.unsysdigital.com/index.php/jias/article/view/10.21535%252Fjias.v3i2.916/0>,
<https://pmc.ncbi.nlm.nih.gov/articles/PMC6389673/>

⁶ <https://pmc.ncbi.nlm.nih.gov/articles/PMC7584400/>, <https://hsph.harvard.edu/exec-ed/news/implementation-gap-artificial-intelligence-help-patients-clinical-practices/>

An OES-style approach would encourage simple, pragmatic evaluation designs—such as rolling pilots in which some clinics receive an AI tool first and others later—so that organizations can compare key outcomes over time without creating extensive research infrastructure. This makes “build evidence while implementing” the norm rather than the exception.

2. An evaluation framework grounded in EHR and EHR-use metadata (Q1, Q2, Q8, Q10)

This section addresses Questions 1, 2, 8, and 10 by proposing a practical framework for measuring AI's effects using data already generated by health systems.

The EHR is already the backbone of digital care delivery. In addition to clinical data, it produces rich metadata about how work is done: who interacts with which records, in what sequence, for how long, and with what patterns of interruption and collaboration. My colleagues and I have shown that this metadata can be organized into “emerging domains” of health care delivery, including team structure and dynamics, workflows and task sequencing, and aspects of the cognitive environment.

For AI, this means we can move beyond counting logins or AI button clicks and ask more meaningful questions, such as:

- Does the AI tool reduce or increase the time spent on routine tasks (e.g., documentation, order entry, message triage)?
- Does it reduce or increase the number of context switches and interruptions in a typical visit?
- Does it change how work is distributed among the care team members?
- Does it reduce after-hours work or shift the burden onto particular roles?

To avoid chasing convenient but low-value metrics, any suggested measures should be screened using three criteria from our EHR-metadata work: ⁷

- Desirable: Does the measure capture something clinicians, patients, or organizations actually care about (for example, less fragmentation, fewer unnecessary clicks, better continuity), rather than simply what is easiest to count?
- Feasible: Can the measure be derived reliably from available data sources, such as EHR metadata or claims, without excessive burden or the need for complex new infrastructure?
- Viable: Can the measure be maintained over time, interpreted correctly, and used for decision-making without creating perverse incentives?

HHS and ONC could use this framework to:

- Curate a small set of recommended AI-related measures for federally supported programs.
- Encourage vendors to expose standardized metadata elements (for example, AI recommendation events, acceptance/override indicators) to support these measures.⁸
- Integrate AI-related metrics into existing EHR reporting and transparency programs, so evaluation becomes part of routine practice rather than an extra task.⁹

3. Closing the “model-to-practice” gap in federally funded AI (Q1, Q2, Q8, Q10)

This section responds to Questions 1, 2, 8, and 10 by focusing on the gap between AI model development and real-world implementation, and how HHS and NIH can better support the full lifecycle.

⁷ <https://www.jmir.org/2025/1/e64721>

⁸ <https://www.hipaajournal.com/onc-proposes-new-rule-to-advance-care-through-technology-and-interoperability/>

⁹ <https://www.ropesgray.com/en/insights/alerts/2023/12/hhs-onc-finalizes-rule-with-new-requirements-for-use-of-ai-in-certified-health-it>

A large share of machine-learning models developed for health and medicine never progress beyond publication. Reviews and commentaries describe an “implementation gap,” where strong model performance on retrospective datasets does not translate into real-world deployment, prospective evaluation, or sustained use in clinical workflows. NIH’s own portfolio analyses show that AI and machine-learning projects now account for a growing share of its research budget. Yet, there is limited visibility into how often these grant-funded models are translated into deployed tools that improve care in practice.¹⁰

In my own work, I have experienced this gap from multiple angles. In a dissertation chapter on causal machine learning for front-of-package food labels to evaluate health food purchases, and in projects on behavioral economics-based machine learning for precision motivation in diabetes prevention, the technical models were only the first step; the hard part was turning them into something that could be implemented, monitored, and iterated in real programs. While the work was presentable at an academic conference, my limitations in academic research training at the time meant I lacked the knowledge to design evaluations or models for operational practices.

To address this, HHS (in coordination with NIH) could:

- Explicitly fund “translation and deployment” phases, not just model development: Create or expand funding mechanisms that require teams to move beyond algorithm development to implementation planning, pilot deployment, and evaluation in real-world settings, ideally with stepwise milestones.
- Require multidisciplinary teams that include operational and implementation expertise: Many academic AI projects are led by scientists and clinicians but lack sustained engagement from operational roles such as implementation engineers, workflow designers, EHR integration specialists, and frontline staff. Funding announcements could set expectations that teams include these perspectives from the outset and that proposals specify how operational partners will be involved in design, testing, and scale-up.
- Tie research funding to basic implementation and evaluation plans: Without turning every grant into a full-scale trial, NIH and other HHS funders could require that AI and machine-learning proposals include a short, concrete plan for: (1) an intended real-world use case, (2) potential integration points (for example, within an EHR or patient-facing tool), and (3) how success and unintended effects would be measured if the model were deployed. This can align with the continuous evaluation expectations described earlier, using EHR and EHR-use metadata where appropriate.
- Create systems of perpetual evaluation: Move from infrequent, “frozen-in-time” evaluations to ongoing, real-time or near-real-time monitoring of AI’s effects, using routine data streams, so that problems are detected and addressed quickly rather than years later in a retrospective study.¹¹

These steps could help shift federally supported AI research from a “model-centric” pipeline that largely ends at publication toward a “deployment-centric” pipeline where development, implementation, and continuous learning are planned as part of a single lifecycle.¹²

4. “Automation of the obvious, flagging of the not-so-obvious” (Q1, Q2, Q9, Q10)

¹⁰ <https://pmc.ncbi.nlm.nih.gov/articles/PMC7584400/>, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12782054/>, <https://pmc.ncbi.nlm.nih.gov/articles/PMC8443295/>, <https://innovations.bmj.com/content/6/2/45>

¹¹ https://acf.gov/sites/default/files/documents/opre/OPRE-MA_Rapid-Learning-Brief_Final-508.pdf,

¹² <https://pmc.ncbi.nlm.nih.gov/articles/PMC8443295/>

This section responds to Questions 1, 2, 9, and 10 by describing how the scope of AI tools can reduce barriers to adoption, support appropriate oversight, and address patient- and caregiver-relevant concerns about safety, attention, and judgment.

AI in health care will be most valuable when it is intentionally scoped. A useful principle is:

- Automate what is routine, predictable, and well-understood (“the obvious”).
- Use AI to highlight patterns, outliers, or risks that are hard for humans to keep track of unaided (“the not-so-obvious”).
- Preserve human expertise and judgment for complex, ambiguous, or sensitive decisions.

In practice, this means:

- Prioritizing AI for tasks like summarizing lengthy histories, retrieving relevant information, standardizing documentation, and handling predictable administrative workflows.
- Using AI to surface subtle patterns (for example, specific combinations of risk factors scattered across notes) but still presenting these as prompts for human review rather than as final decisions.
- Avoiding models that try to replace nuanced clinical judgment in high-stakes settings without strong evidence and stringent guardrails.

Evaluation should explicitly test whether AI is achieving this balance. For example, EHR metadata can show whether documentation time decreases without an offsetting increase in corrections or follow-up, or whether certain types of AI-driven alerts lead to better, not just more, interventions.

5. Incentives that reward measured value, not mere adoption (Q1, Q2, Q10)

This section addresses Questions 1, 2, and 10 by suggesting how HHS payment and grant structures can reduce barriers to effective AI use (Q1), align program design with meaningful measures (Q2), and prioritize evaluation that focuses on value rather than volume (Q10).

Payment and grant policies will shape how and where AI is used. If incentives reward adoption alone, organizations may deploy tools that add little value or even increase burden.

Instead, HHS can:

- Link support to improvements in selected measures: Make AI-related funding conditional on demonstrating meaningful improvement on a small set of agreed-upon, context-appropriate measures. These could include reduced after-hours work, shorter time to complete key workflows, improved coordination, or timelier follow-up, measured using EHR metadata and related data sources.
- Allow tailored measure selection within guardrails: Recognize that different settings (for example, a rural clinic vs a large academic center) will prioritize different outcomes. Within the desirability/feasibility/viability framework, allow organizations to choose metric bundles that fit their context, provided they are clearly tied to patient care and staff well-being.
- Avoid tying incentives to “AI usage” alone: Avoid metrics like “number of AI-assisted encounters” as primary targets, since these can be increased without improving care and can invite overuse.

This approach keeps the focus on value and outcomes while giving organizations flexibility in how they get there.

6. Guarding against automation complacency and overload (Q1, Q2, Q9, Q10)

This section addresses Questions 1, 2, 9, and 10 by highlighting new safety barriers AI may introduce, suggesting regulatory expectations for human-centered design and training, and identifying areas for ongoing safety research.

Lessons from other sectors, such as aviation, show that automation can improve safety and efficiency while also creating new risks, including over-reliance and the erosion of manual skills. Health care is vulnerable to similar patterns, in which training and apprenticeship can rely heavily on AI-based assistance, sacrificing the effort needed to develop skilled, knowledgeable health care workers.

To address this, HHS could:

- Expect human-centered design: Encourage or require that AI tools in clinical care be designed to keep humans appropriately engaged. This includes clear presentation of uncertainty, easy ways to override or decline AI recommendations, and transparency about when AI is active or inactive.¹³
- Monitor for over-reliance and new error patterns: Use EHR- and metadata-derived measures (e.g., changes in override rates, new clusters of near-misses, or unusual patterns of documentation or task switching) as early warning signs of problematic reliance on automation.
- Support training that anticipates AI failures: Promote training and simulation that includes scenarios where AI is wrong, delayed, or offline, so clinicians and staff practice operating under these conditions. The goal is to ensure they can recognize when something is off and respond appropriately, as other high-reliability sectors train for automation failures.

Integrating these expectations into federal programs and guidance will help ensure that AI augments, rather than displaces, human expertise and situational awareness.

7. Supporting adoption in resource-constrained and non-tech-hub settings (Q1, Q2, Q8, Q9)

This section responds to Questions 1, 2, 8, and 9 by describing barriers that rural, small, and resource-constrained organizations face, suggesting program design and data-collection approaches that make AI adoption more feasible across settings, and addressing patient and caregiver concerns about uneven access to AI-enabled services.

Experience with telemedicine has shown that new technologies often reach well-resourced, urban, or academic systems first, and only later (if at all) reach small practices, rural facilities, and other resource-constrained settings. AI is likely to follow the same pattern unless policy explicitly pushes against it.¹⁴

To avoid deepening gaps between organizations, HHS can:

- Pair AI incentives with infrastructure and support: When offering AI-related funding, set aside resources for technical assistance, training, and infrastructure in rural areas, small practices, and safety-net organizations, including those serving veterans and other high-need populations.
- Design for low-resource environments: Encourage AI solutions that can function in lower-bandwidth, less tightly integrated IT environments, rather than assuming the capabilities of large, integrated delivery systems.
- Require reporting by organization type and setting. For federally supported AI implementations, request that evaluation results be stratified by organization type (e.g., small vs. large practice) and geography (e.g., rural vs. urban). This stratification will help HHS detect and address emerging “AI haves and have-nots” among organizations, not just individuals.

These actions align AI investments with the goal of broadly distributed benefits, rather than concentrating them in areas where infrastructure is already strong.

¹³ <https://pmc.ncbi.nlm.nih.gov/articles/PMC9557254/>

¹⁴ <https://pmc.ncbi.nlm.nih.gov/articles/PMC9719760/>,
<https://journals.plos.org/digitalhealth/article?id=10.1371%2Fjournal.pdig.0000952>

Conclusion

AI in health care will not succeed on technical performance alone. It must be implementable in real-world settings, usable without adding burden, and continuously monitored to ensure it delivers intended benefits and does not create new problems. EHR metadata and EHR use provide a powerful, underused lens for understanding how AI changes workflows, team dynamics, and cognitive environments. Telemedicine's uneven diffusion has shown how easily digital innovations can widen gaps between organizations and communities when implementation is not intentional. By requiring continuous evaluation, grounding measures in desirable, feasible, and viable metrics, rewarding demonstrated value rather than mere adoption, closing the model-to-practice gap, and intentionally supporting resource-constrained settings, HHS can guide AI from promise to practical impact.

Thank you for considering these comments. I would be glad to provide additional details on any of the points above.

Sincerely,

Ratnalekha Viswanadham, Ph.D.

<https://www.linkedin.com/in/dr-ratnalekha-viswanadham/>